

<http://poloclub.gatech.edu/cse6242>

CSE6242 / CX4242: Data & Visual Analytics

# Text Analytics (Text Mining)

Concepts, Algorithms, LSI/SVD

Duen Horng (Polo) Chau

Assistant Professor

Associate Director, MS Analytics

Georgia Tech

Partly based on materials by

Professors Guy Lebanon, Jeffrey Heer, John Stasko, Christos Faloutsos, Parishit Ram (GT PhD alum; SkyTree), Alex Gray

# Text is everywhere

We use documents as primary information artifact in our lives

Our access to documents has grown tremendously thanks to the Internet

- *WWW*: webpages, Twitter, Facebook, Wikipedia, Blogs, ...
- *Digital libraries*: Google books, ACM, IEEE, ...
- Lyrics, closed caption... (youtube)
- Police case reports
- Legislation (law)
- Reviews (products, rotten tomatoes)
- Medical reports (EHR - electronic health records)
- Job descriptions

# Big (Research) Questions

... in understanding and gathering information from text and document collections

- establish authorship, authenticity; plagiarism detection
- classification of genres for narratives (e.g., books, articles)
- tone classification; sentiment analysis (online reviews, twitter, social media)
- code: syntax analysis (e.g., find common bugs from students' answers)

# Popular Natural Language Processing (NLP) libraries

- **Stanford NLP**

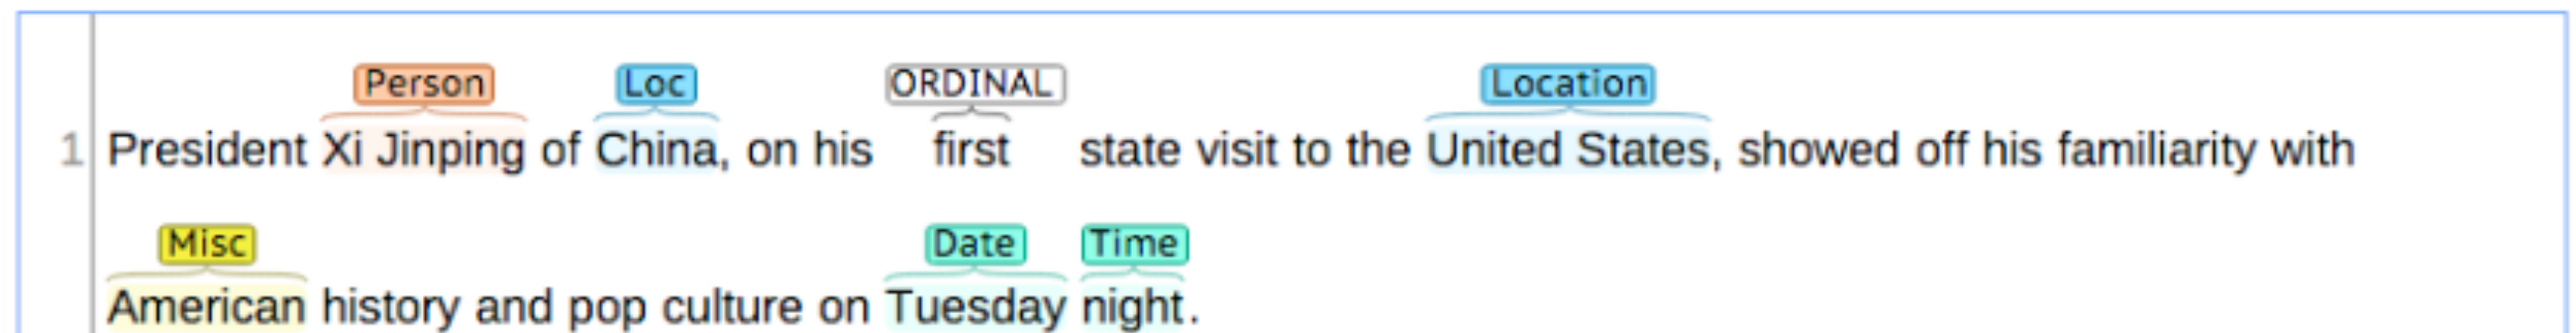
- **OpenNLP**

- **NLTK (python)**

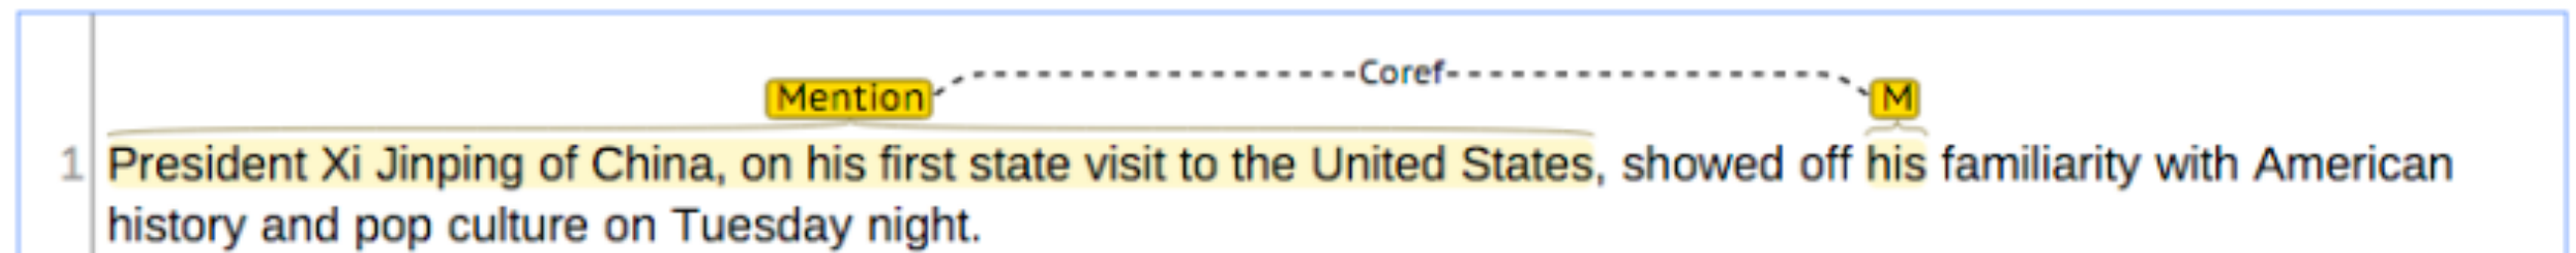
tokenization, sentence segmentation, part-of-speech tagging, named entity extraction, chunking, parsing

## Named Entity Recognition:

Image source: <https://stanfordnlp.github.io/CoreNLP/>



## Coreference:



## Basic Dependencies:

# Outline

- **Preprocessing** (e.g., stemming, remove stop words)
- **Document representation** (most common: bag-of-words model)
- **Word importance** (e.g., word count, TF-IDF)
- **Latent Semantic Indexing** (find “concepts” among documents and words), which helps with **retrieval**

To learn more: Prof. Jacob Eisenstein's  
**CS 4650/7650 Natural Language Processing**

# Stemming

Reduce words to their **stems** (or base forms)

**Words:** compute, computing, computer, ...

**Stem:** comput

Several classes of algorithms to do this:

- Stripping suffixes, lookup-based, etc.

<http://en.wikipedia.org/wiki/Stemming>

Stop words: [http://en.wikipedia.org/wiki/Stop\\_words](http://en.wikipedia.org/wiki/Stop_words)

# Bag-of-words model

Represent each **document** as a **bag of words**, ignoring words' ordering. Why? For **simplicity**.

Unstructured text becomes **a vector of numbers**

e.g., docs: “I like visualization”, “I like data”.

1 : “I”

2 : “like”

3 : “data”

4 : “visualization”

“I like visualization”  $\Rightarrow$  [1, 1, 0, 1]

“I like data”  $\Rightarrow$  [1, 1, 1, 0]

# TF-IDF

A word's importance score in a document, among  $N$  documents

**When** to use it? Everywhere you use “word count”, you can likely use TF-IDF.

**TF:** term frequency

= #appearance a document

(high, if terms appear many times in this document)

**IDF:** inverse document frequency

=  $\log( N / \text{\#document containing that term} )$

(penalize “common” words appearing in almost any documents)

**Final score = TF \* IDF**

(higher score  $\Rightarrow$  more “characteristic”)



# Vector Space Model

## Why?

Each document  $\Rightarrow$  vector

Each query  $\Rightarrow$  vector

Search for documents  $\Rightarrow$  find “similar” vectors

Cluster documents  $\Rightarrow$  cluster “similar” vectors

# Latent Semantic Indexing (LSI)

Main idea

- map each **document** into some ‘**concepts**’
- map each **term** into some ‘**concepts**’

‘**Concept**’ : ~ a set of terms, with weights.

For example, **DBMS\_concept**:

“data” (0.8),

“system” (0.5),

# Latent Semantic Indexing (LSI)

*~ pictorially (before) ~*

**document-term** matrix

	data	system	retireval	lung	ear
doc1	1	1	1		
doc2	1	1	1		
doc3				1	1
doc4				1	1

# Latent Semantic Indexing (LSI)

*~ pictorially (after) ~*

**term-concept**  
matrix

	database concept	medical concept
data	1	
system	1	
retrieval	1	
lung		1
ear		1

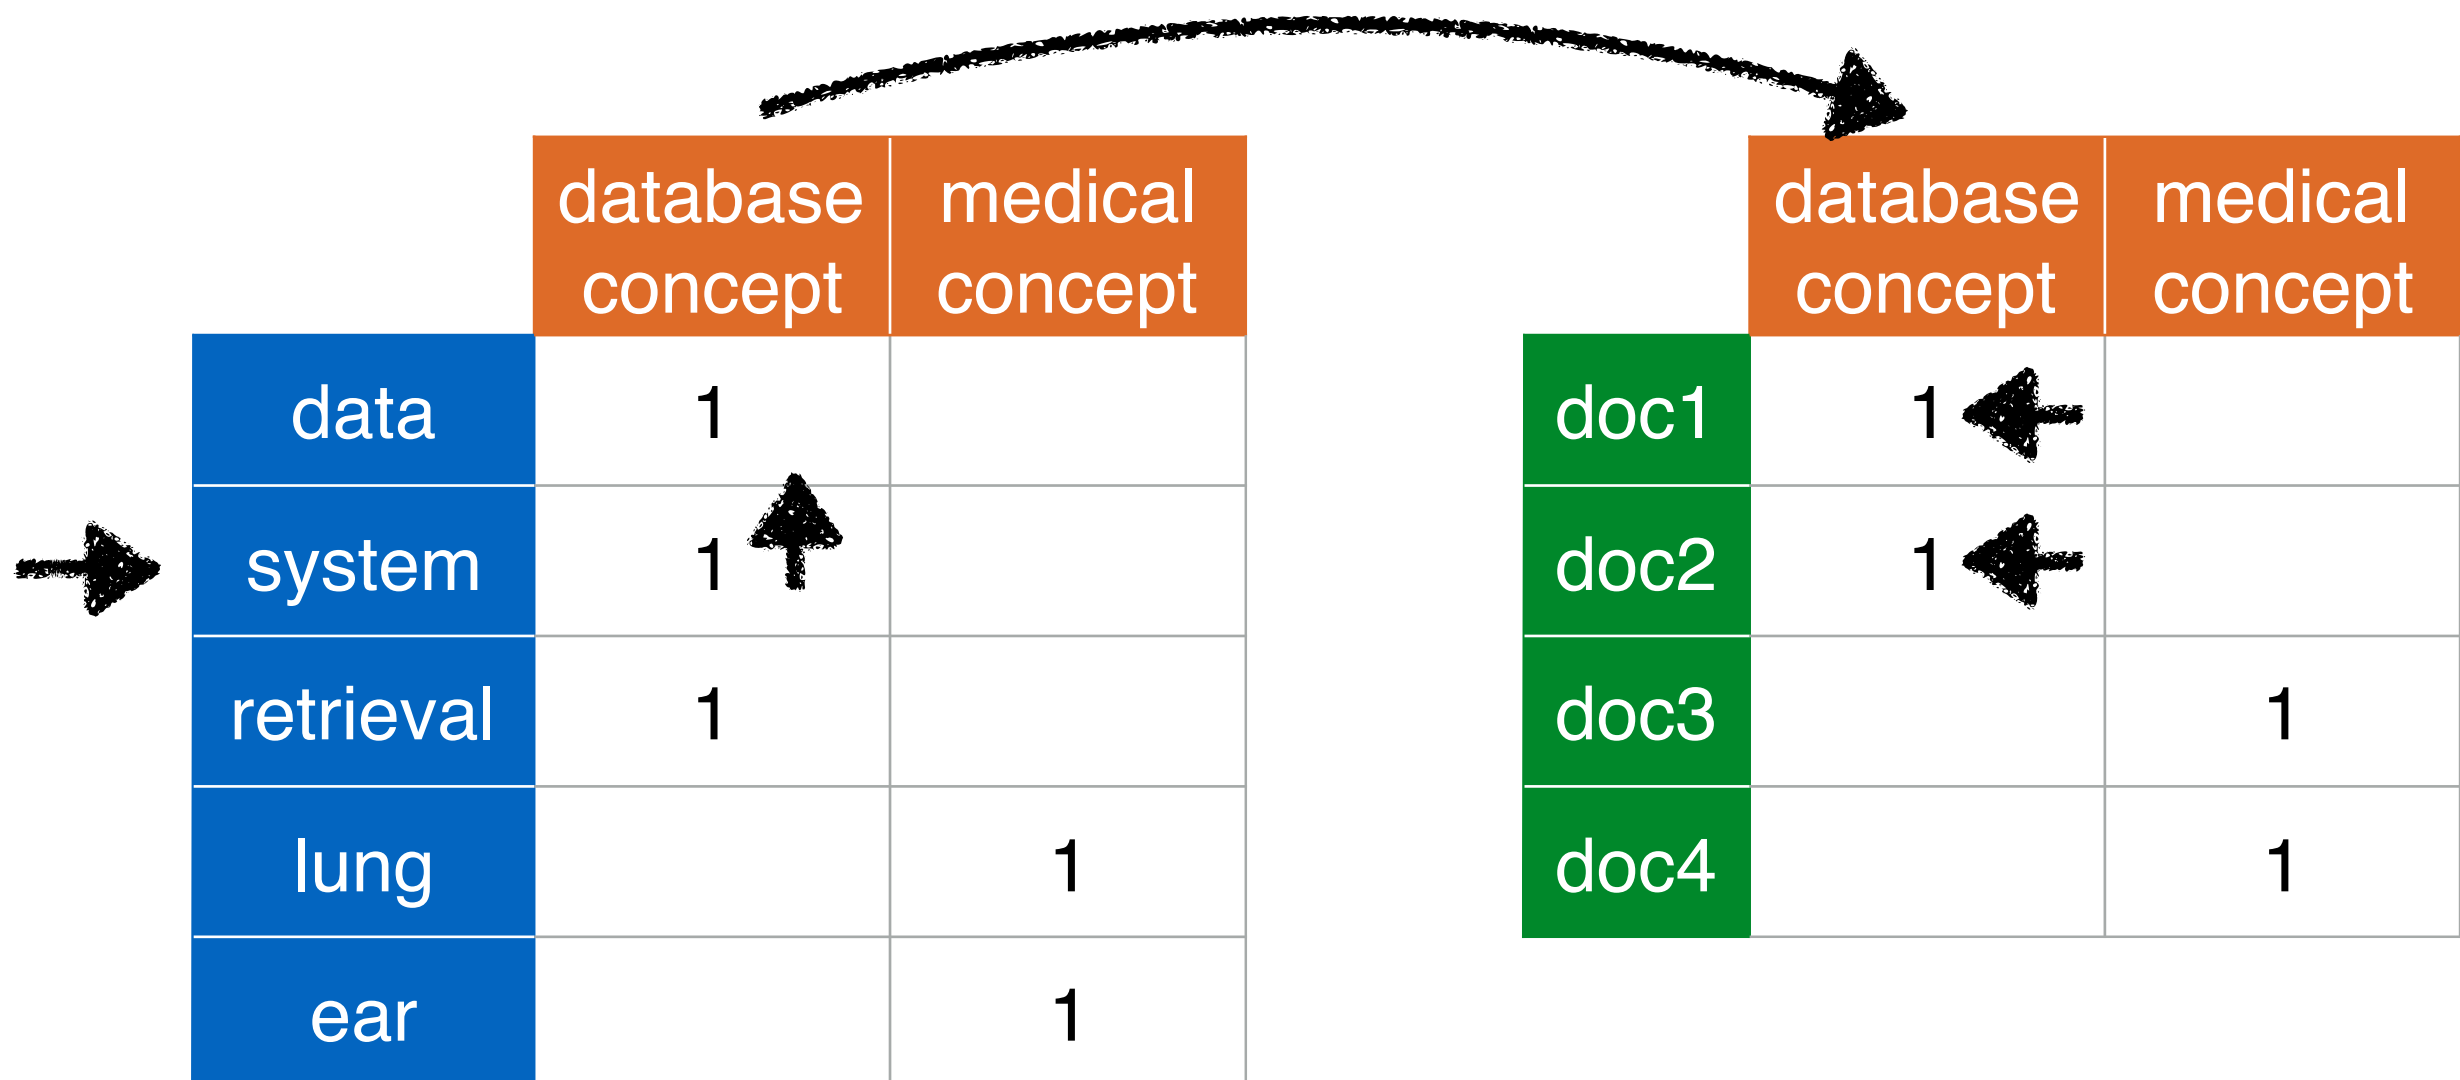
*... and*  
**document-concept**  
matrix

	database concept	medical concept
doc1	1	
doc2	1	
doc3		1
doc4		1

# Latent Semantic Indexing (LSI)

Q: How to search, e.g., for “system”?

A: find the corresponding concept(s); and the corresponding documents



# Latent Semantic Indexing (LSI)

Works like an **automatically constructed thesaurus**

We may retrieve documents that **DON'T** have the term “system”, but they contain almost everything else (“data”, “retrieval”)

# LSI - Discussion

Great idea,

- to derive ‘**concepts**’ from documents
- to build a ‘**thesaurus**’ automatically
- to reduce dimensionality (down to few “concepts”)

How does LSI work?

Uses **Singular Value Decomposition** (SVD)

# Singular Value Decomposition (SVD)

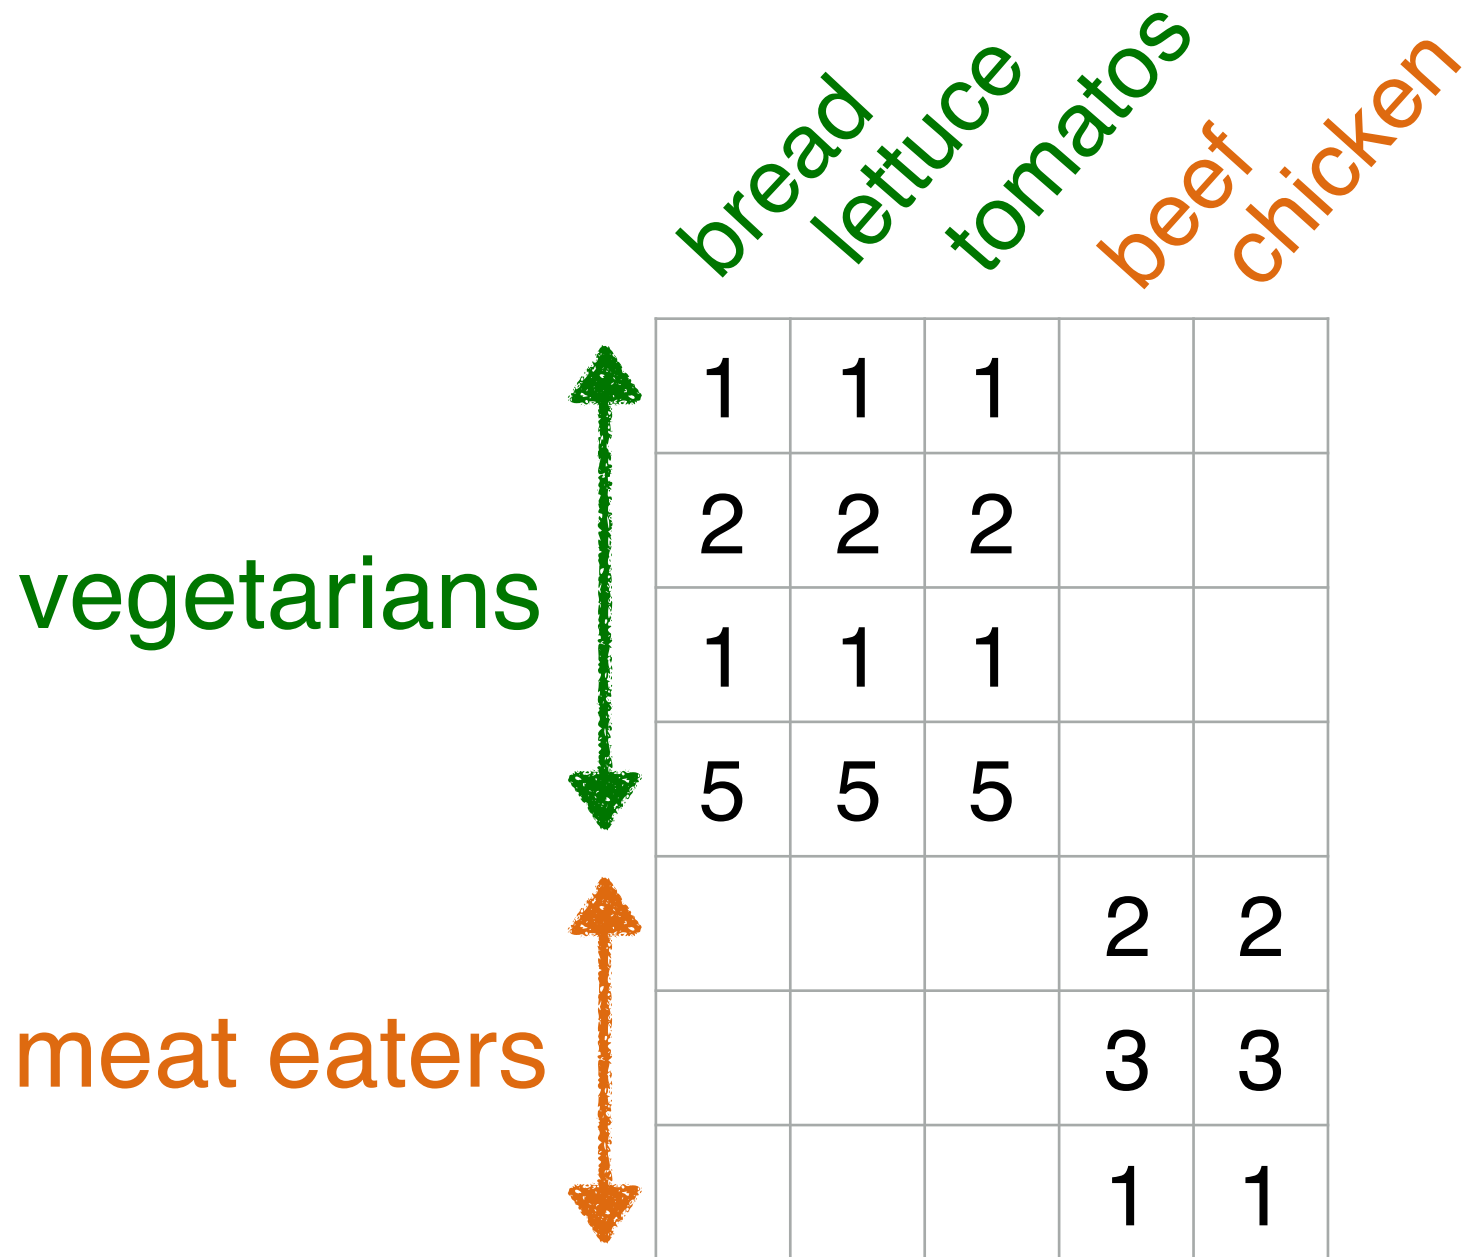
## Motivation

### Problem #1

Find “concepts”  
in matrices

### Problem #2

Compression /  
dimensionality  
reduction



	bread	lettuce	tomatos	beef	chicken
vegetarians	1	1	1		
	2	2	2		
	1	1	1		
	5	5	5		
meat eaters				2	2
				3	3
				1	1



# SVD is a **powerful,** **generalizable** technique.

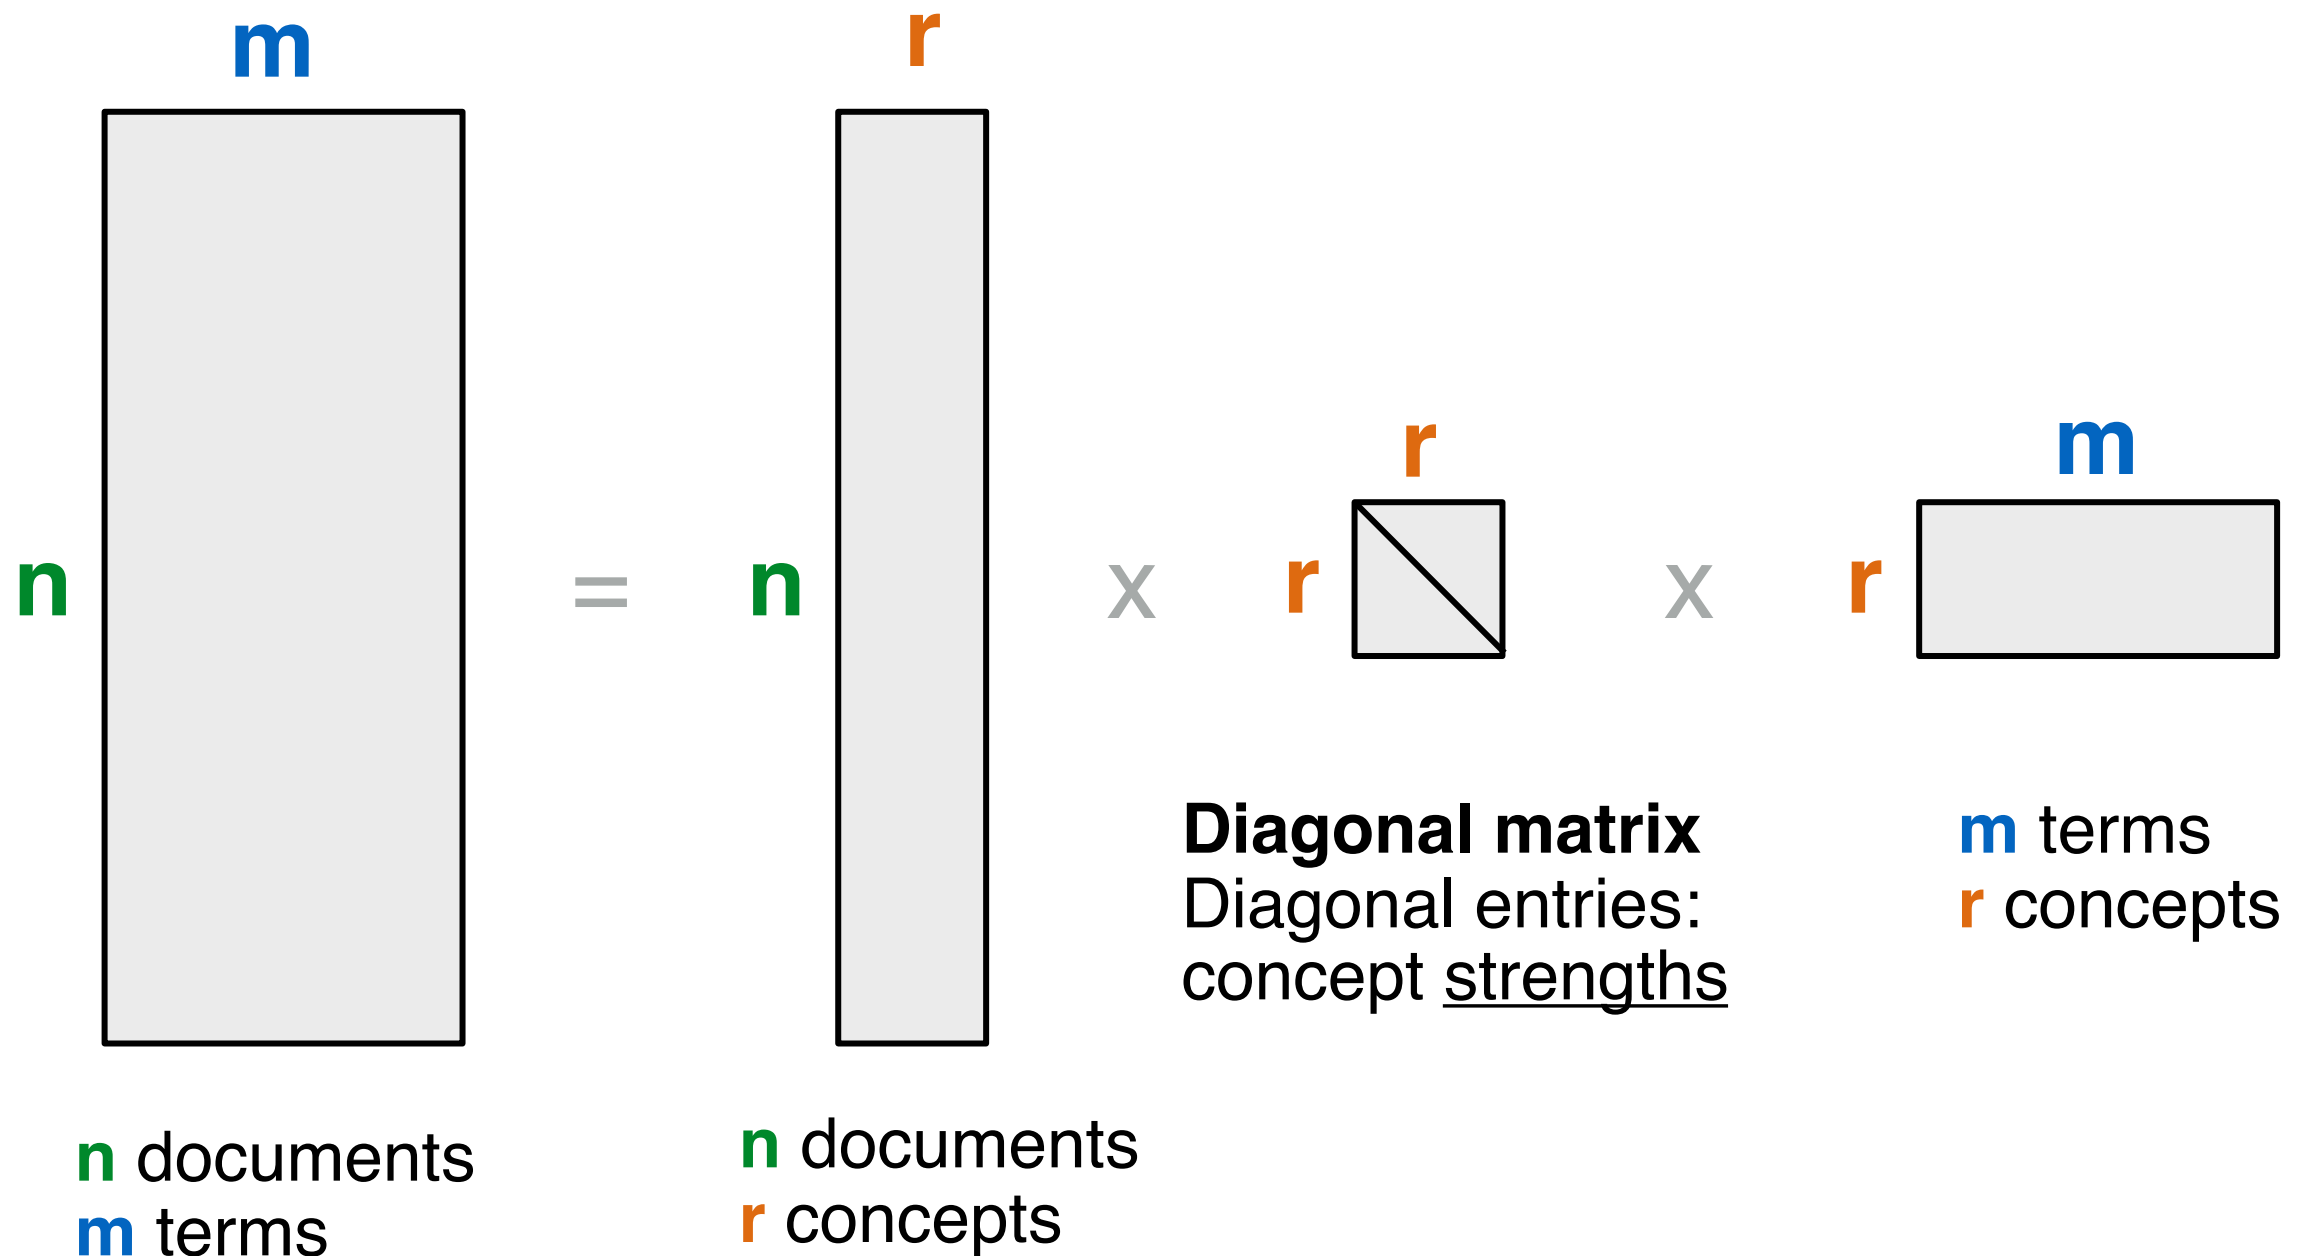
Songs / Movies / Products

Customers

1	1	1		
2	2	2		
1	1	1		
5	5	5		
			2	2
			3	3
			1	1

# SVD Definition (pictorially)

$$\mathbf{A}_{[n \times m]} = \mathbf{U}_{[n \times r]} \mathbf{\Lambda}_{[r \times r]} (\mathbf{V}_{[m \times r]})^T$$



# SVD Definition (in words)

$$\mathbf{A}_{[n \times m]} = \mathbf{U}_{[n \times r]} \mathbf{\Lambda}_{[r \times r]} (\mathbf{V}_{[m \times r]})^T$$

**A: n x m matrix**

e.g., n documents, m terms

**U: n x r matrix**


e.g., n documents, r concepts

**$\mathbf{\Lambda}$ : r x r diagonal matrix**

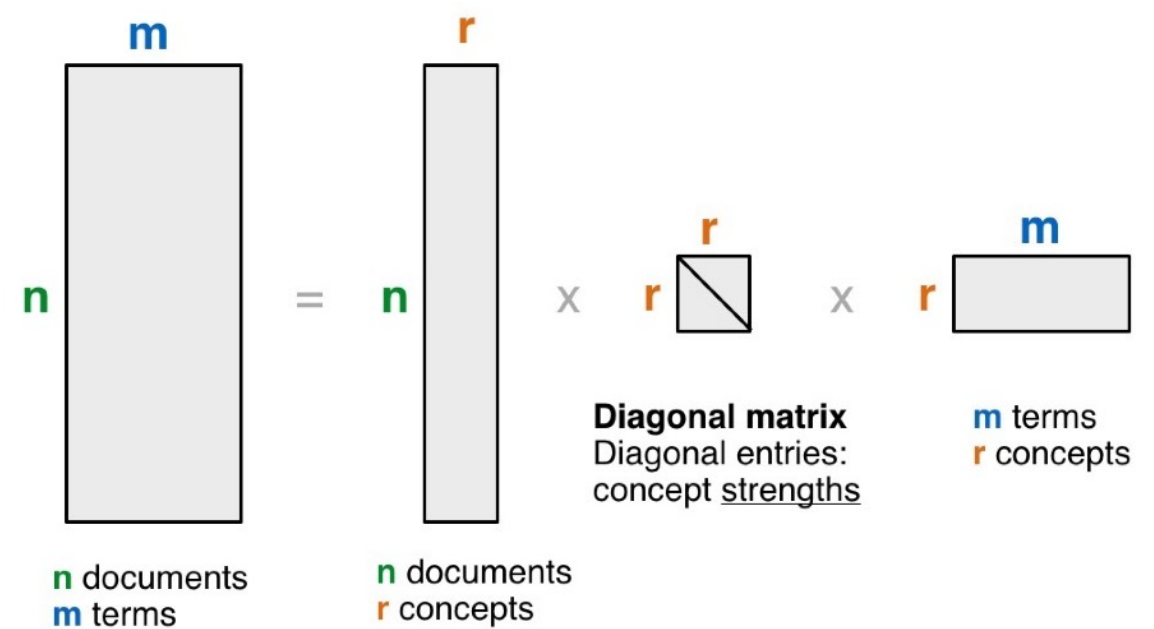
r : rank of the matrix; strength of each 'concept'

**V: m x r matrix**

e.g., m terms, r concepts



# SVD - Properties



**THEOREM [Press+92]:**

**always possible to decompose** matrix **A** into

$$A = U \Lambda V^T$$

**U,  $\Lambda$ , V:** **unique**, most of the time

**U, V:** column **orthonormal**


i.e., columns are **unit vectors**, and **orthogonal** to each other

$$U^T U = I$$

$$V^T V = I \quad (I: \text{identity matrix})$$

**$\Lambda$ :** **diagonal** matrix with non-negative diagonal entries, sorted in **decreasing order**

# SVD - Example



	data	info	retrieval	brain	lung
CS docs	1	1	1	0	0
	2	2	2	0	0
	1	1	1	0	0
	5	5	5	0	0
MD docs	0	0	0	2	2
	0	0	0	3	3
	0	0	0	1	1

=

0.18	0
0.36	0
0.18	0
0.90	0
0	0.53
0	0.80
0	0.27

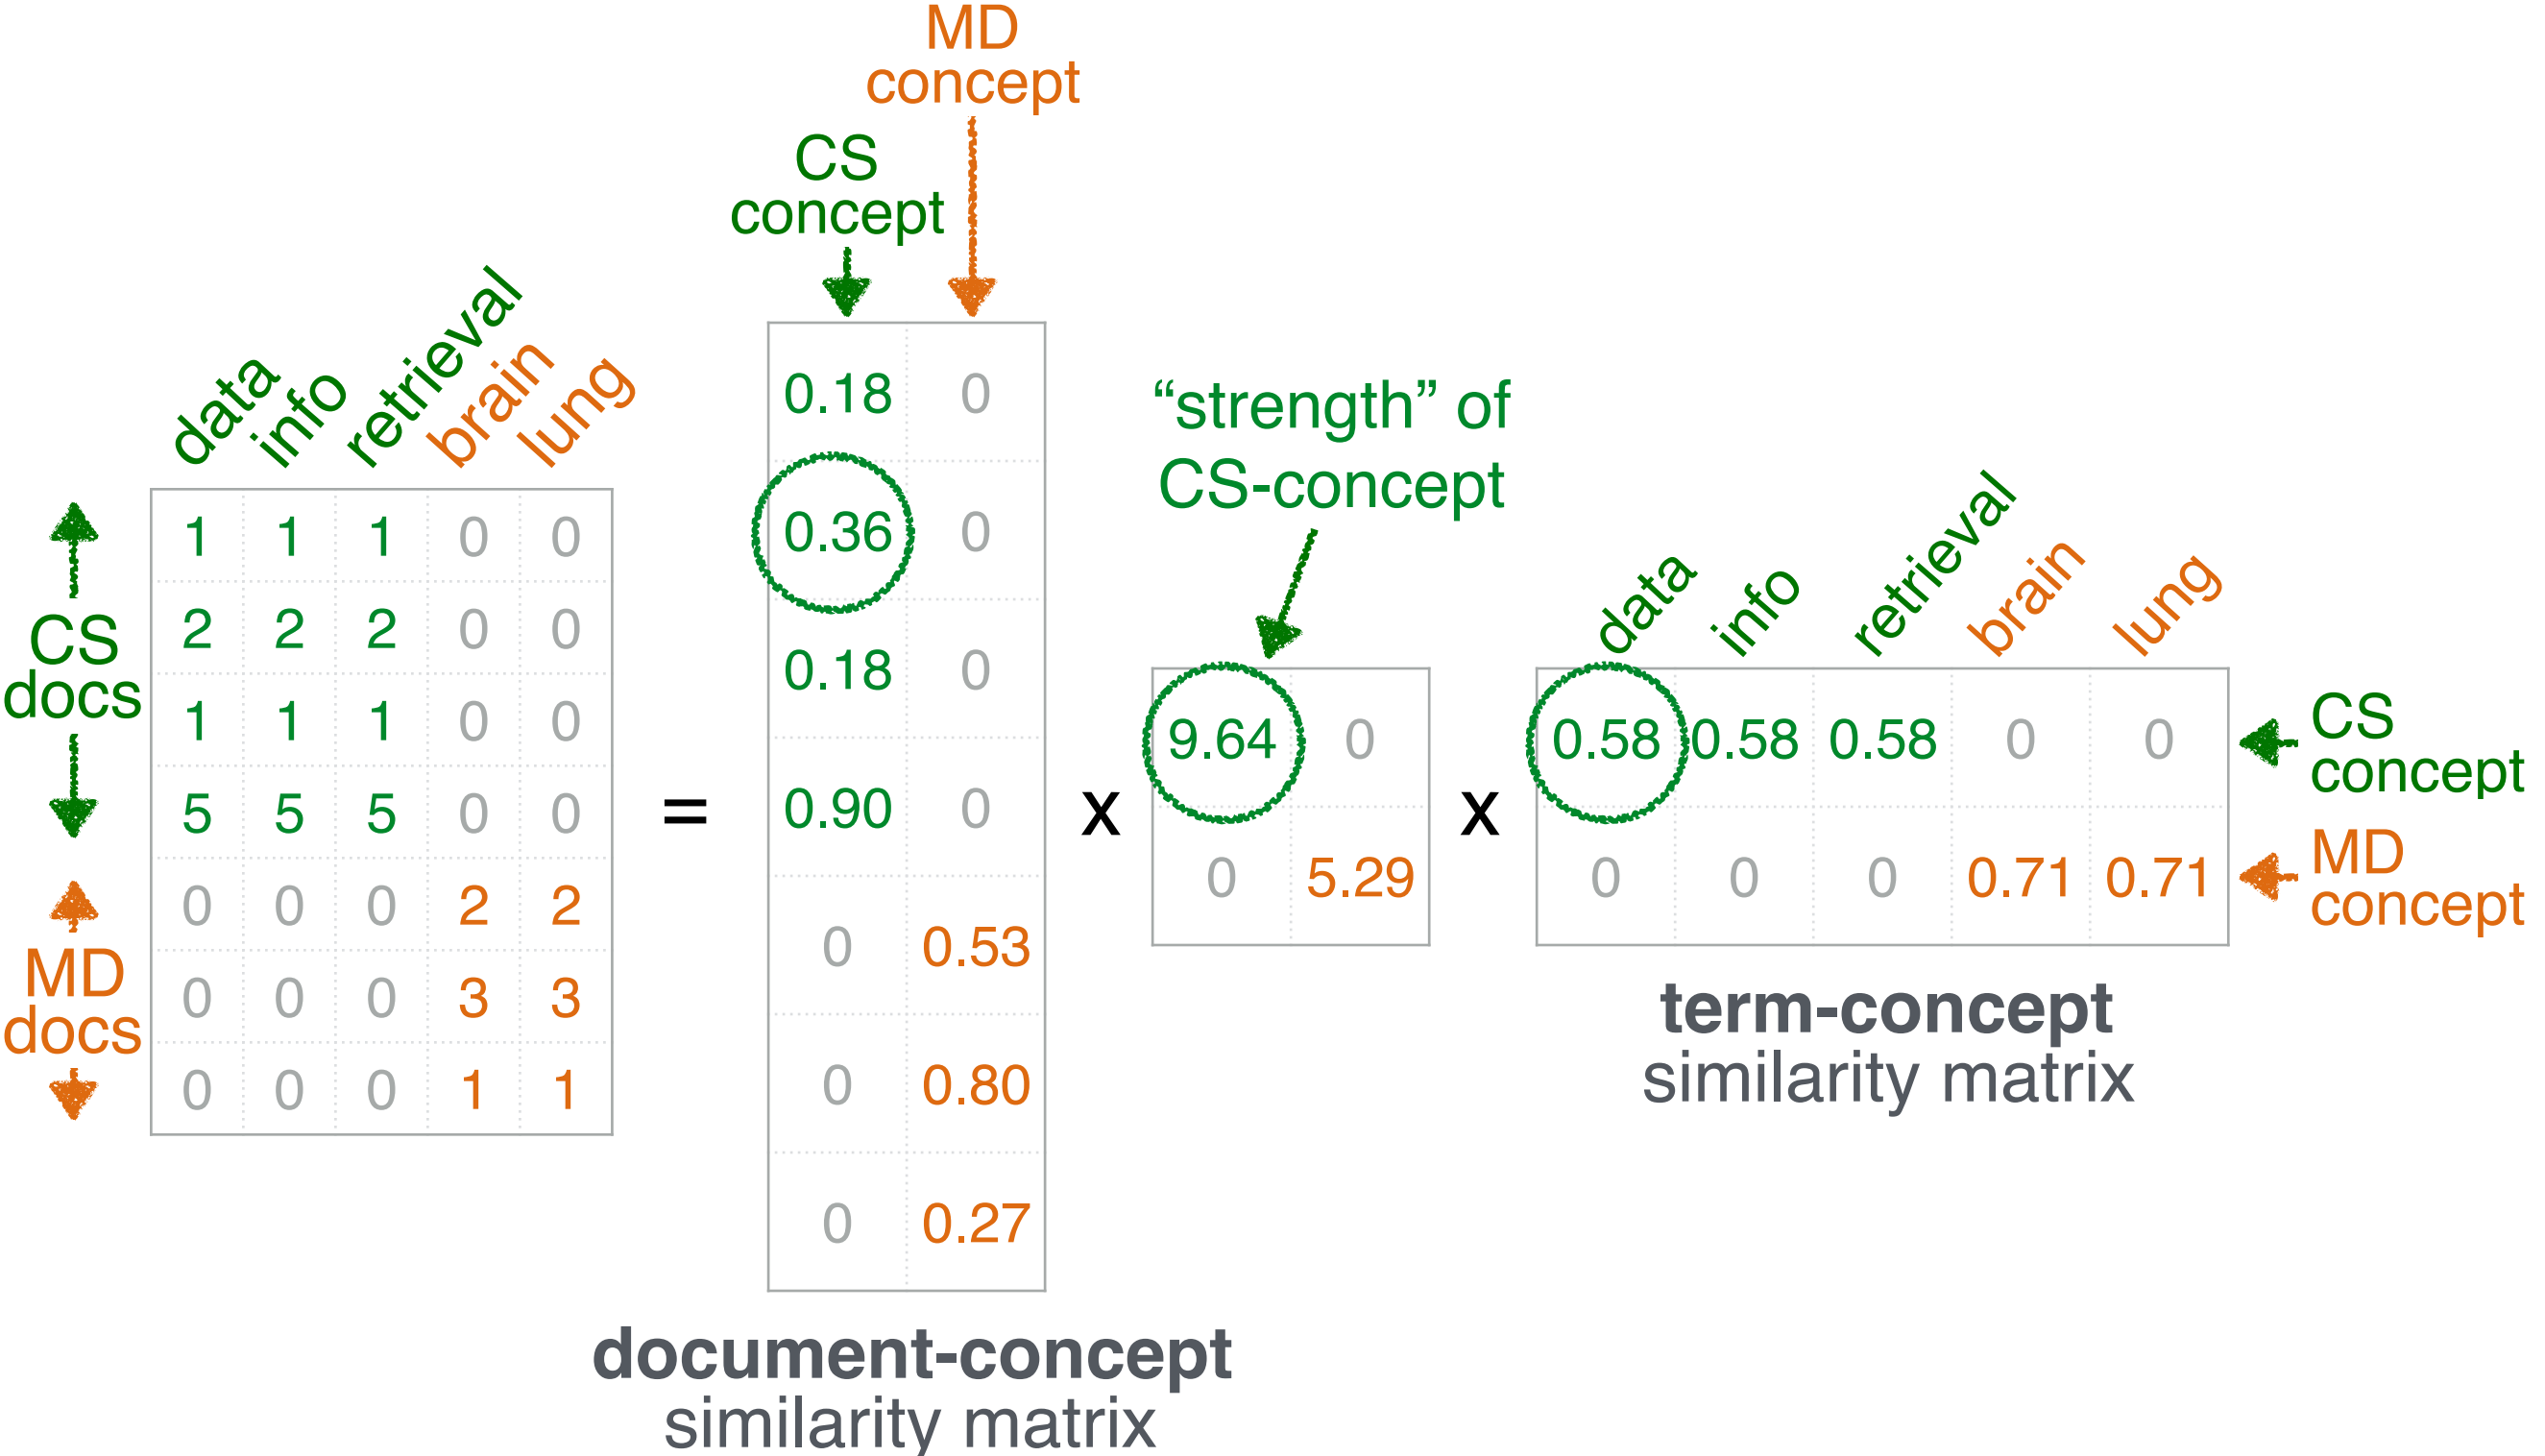
x

9.64	0
0	5.29

x

0.58	0.58	0.58	0	0
0	0	0	0.71	0.71

# SVD - Example



# SVD - Interpretation #1

‘documents’, ‘terms’ and ‘concepts’:

**U**: document-concept similarity matrix

**V**: term-concept similarity matrix

**$\Lambda$** : diagonal elements: concept “strengths”

# SVD - Interpretation #1

‘documents’, ‘terms’ and ‘concepts’:

Q: if  $\mathbf{A}$  is the document-to-term matrix,  
what is the similarity matrix  $\mathbf{A}^T \mathbf{A}$  ?

A:

Q:  $\mathbf{A} \mathbf{A}^T$  ?

A:



# SVD - Interpretation #1

‘documents’, ‘terms’ and ‘concepts’:

Q: if  $\mathbf{A}$  is the document-to-term matrix,  
what is the similarity matrix  $\mathbf{A}^T \mathbf{A}$  ?

A: term-to-term ( $[m \times m]$ ) similarity matrix

Q:  $\mathbf{A} \mathbf{A}^T$  ?

A: document-to-document ( $[n \times n]$ ) similarity matrix

# SVD properties

**V** are the eigenvectors of the *covariance matrix*  $\mathbf{A}^T\mathbf{A}$

$$\mathbf{A}^T\mathbf{A} = (\mathbf{U}\Sigma\mathbf{V}^T)^T (\mathbf{U}\Sigma\mathbf{V}^T) = \mathbf{V}\Sigma^2\mathbf{V}^T$$

**U** are the eigenvectors of the *Gram (inner-product) matrix*  $\mathbf{A}\mathbf{A}^T$

$$\mathbf{A}\mathbf{A}^T = (\mathbf{U}\Sigma\mathbf{V}^T)(\mathbf{U}\Sigma\mathbf{V}^T)^T = \mathbf{U}\Sigma^2\mathbf{U}^T$$

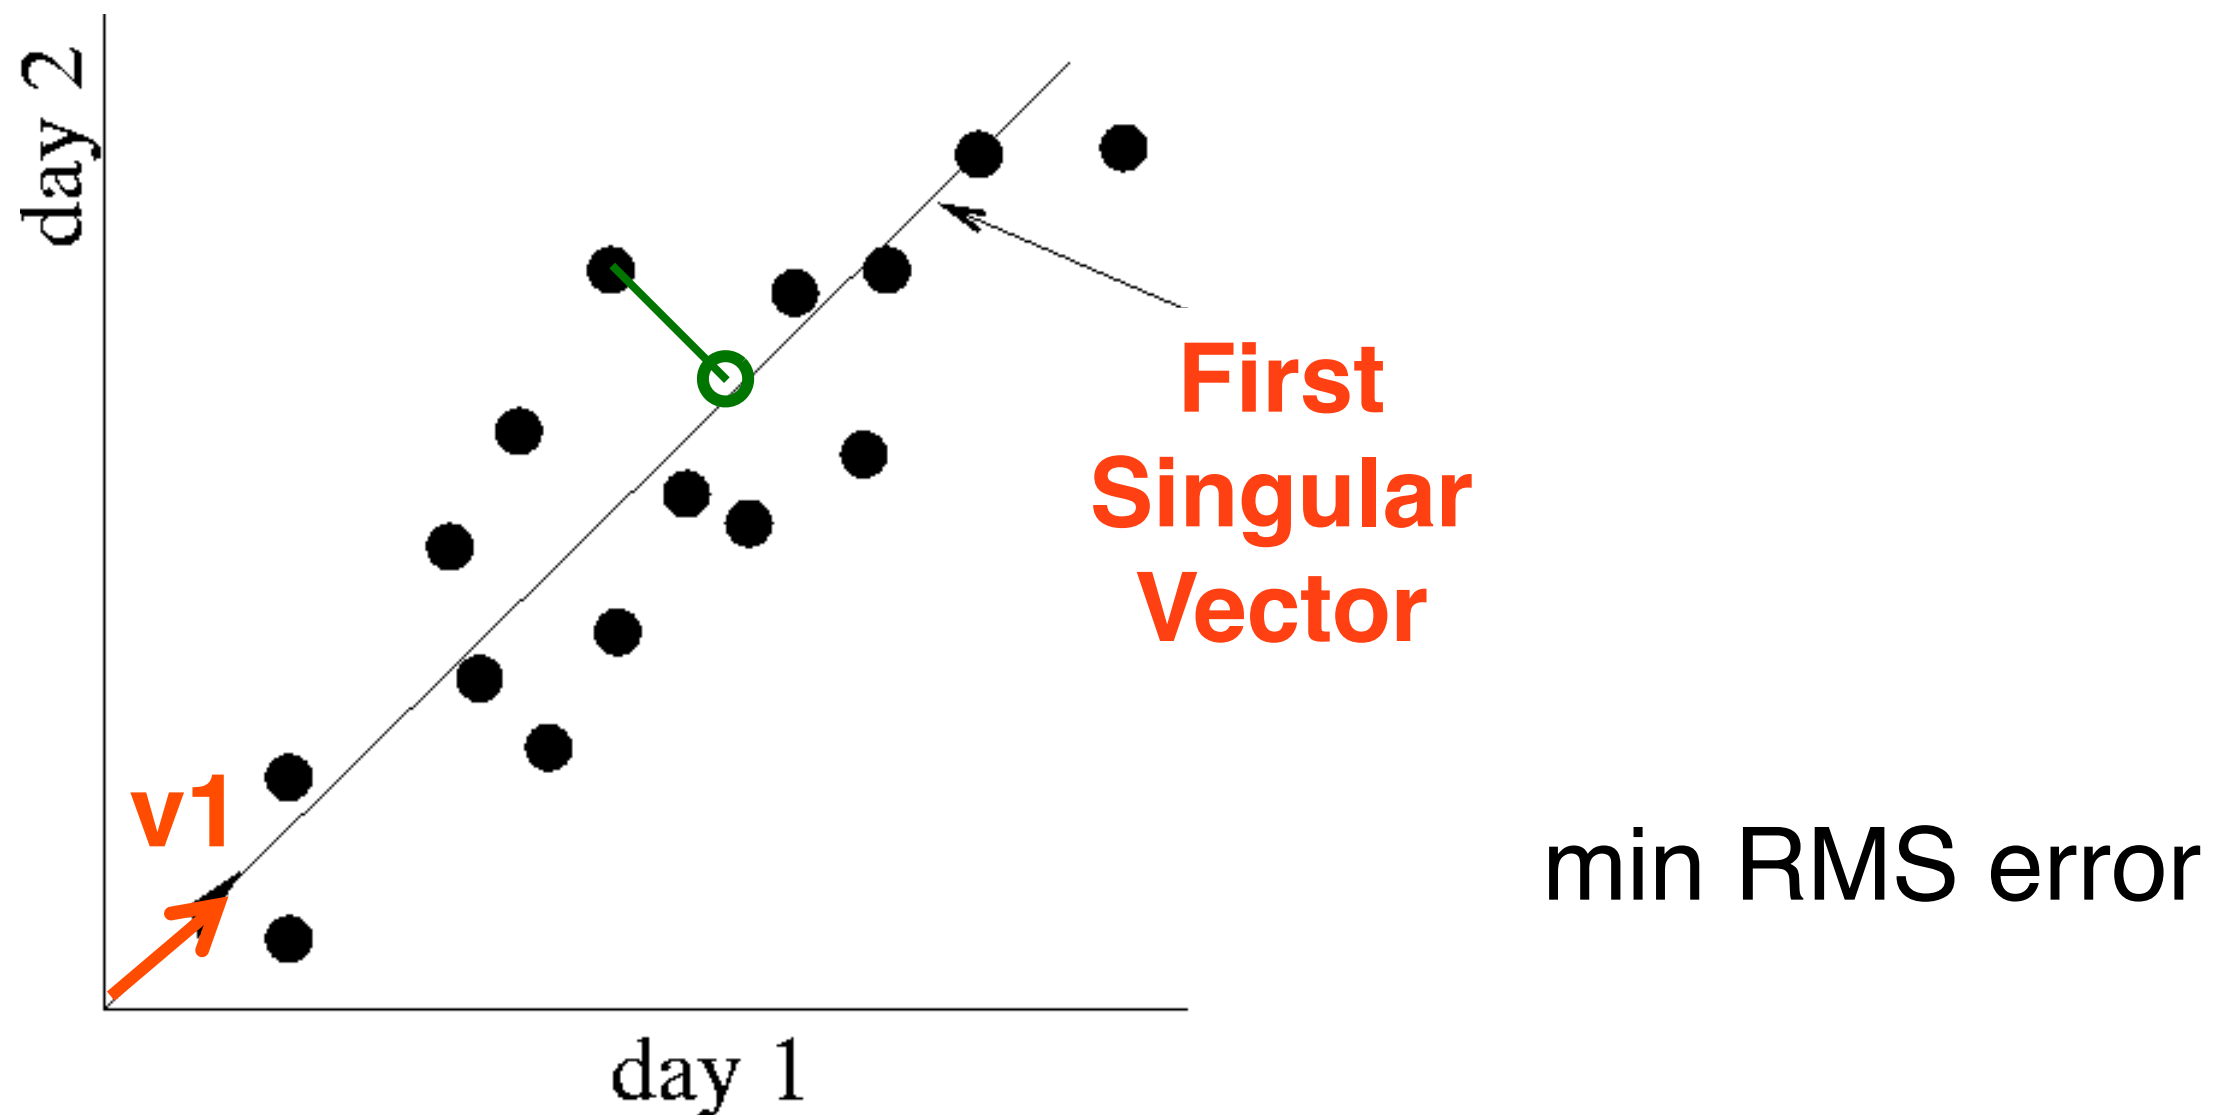
SVD is closely related to PCA, and can be numerically more stable.  
For more info, see:

<http://math.stackexchange.com/questions/3869/what-is-the-intuitive-relationship-between-svd-and-pca> Ian T. Jolliffe, Principal Component Analysis (2nd ed), Springer, 2002. Gilbert Strang, Linear Algebra and Its Applications (4th ed), Brooks Cole, 2005.

# SVD - Interpretation #2

**Find the best axis to project on.**


(‘best’ = min sum of squares of projection errors)



Beautiful visualization explaining PCA:  
<http://setosa.io/ev/principal-component-analysis/>

# SVD - Interpretation #2

$$\mathbf{A} = \mathbf{U} \mathbf{\Lambda} \mathbf{V}^T$$



1	1	1	0	0
2	2	2	0	0
1	1	1	0	0
5	5	5	0	0
0	0	0	2	2
0	0	0	3	3
0	0	0	1	1

=

0.18	0
0.36	0
0.18	0
0.90	0
0	0.53
0	0.80
0	0.27

x

9.64	0
0	5.29

x

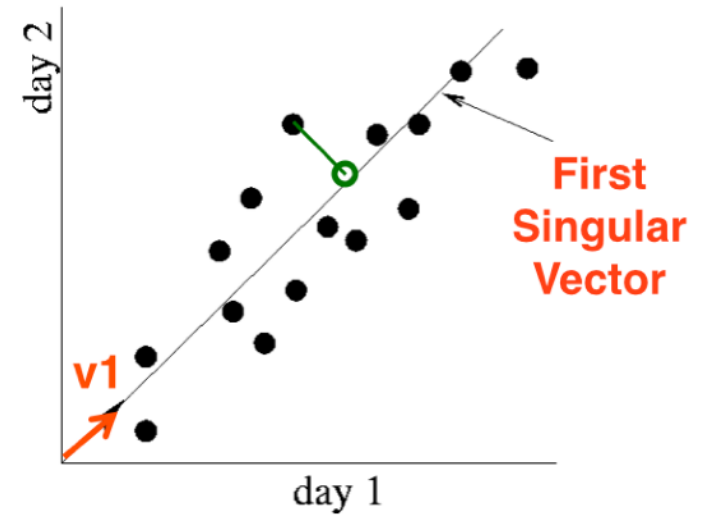
0.58	0.58	0.58	0	0
0	0	0	0.71	0.71

variance ('spread')  
on the v1 axis

v1

# SVD - Interpretation #2

$U \Lambda$  gives the **coordinates** of the points in the projection axis



$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix} \times \begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix} \times \begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$$

$v1$

# SVD - Interpretation #2

More details

Q: how exactly is dim. reduction done?

1	1	1	0	0
2	2	2	0	0
1	1	1	0	0
5	5	5	0	0
0	0	0	2	2
0	0	0	3	3
0	0	0	1	1

=

0.18	0
0.36	0
0.18	0
0.90	0
0	0.53
0	0.80
0	0.27

x

9.64	0
0	5.29

x

0.58	0.58	0.58	0	0
0	0	0	0.71	0.71

# SVD - Interpretation #2

More details

Q: how exactly is dim. reduction done?

A: set the smallest singular values to zero:

1	1	1	0	0
2	2	2	0	0
1	1	1	0	0
5	5	5	0	0
0	0	0	2	2
0	0	0	3	3
0	0	0	1	1

=

0.18	0
0.36	0
0.18	0
0.90	0
0	0.53
0	0.80
0	0.27

x

9.64	0
0	<del>5.19</del>

x

0.58	0.58	0.58	0	0
0	0	0	0.71	0.71

# SVD - Interpretation #2

More details

Q: how exactly is dim. reduction done?

A: set the smallest singular values to zero:

The diagram illustrates the process of dimensionality reduction in SVD. It shows the decomposition of a matrix into three components:  $U$ ,  $\Sigma$ , and  $V^T$ . The matrix  $U$  is a 7x5 matrix with columns 1, 2, and 3 highlighted in green and columns 4 and 5 in orange. The matrix  $\Sigma$  is a 7x2 matrix with the first column highlighted in green and the second in orange. The matrix  $V^T$  is a 2x5 matrix with the first three columns highlighted in green and the last two in orange. Red boxes and 'X' marks indicate that the smallest singular values (0.53, 0.80, and 0.27) are being set to zero to reduce the dimensionality.

1	1	1	0	0
2	2	2	0	0
1	1	1	0	0
5	5	5	0	0
0	0	0	2	2
0	0	0	3	3
0	0	0	1	1

=

0.18	0
0.36	0
0.18	0
0.90	0
0	0.53
0	0.80
0	0.27

x

9.64	0
0	5.19

x

0.58	0.58	0.58	0	0
0	0	0	0.71	0.71



# SVD - Interpretation #2

More details

Q: how exactly is dim. reduction done?

A: set the smallest singular values to zero:

1	1	1	0	0
2	2	2	0	0
1	1	1	0	0
5	5	5	0	0
0	0	0	2	2
0	0	0	3	3
0	0	0	1	1

=

0.18	
0.36	
0.18	
0.90	
0	
0	
0	

x

9.64	0
0	

x

0.58	0.58	0.58	0	0

# SVD - Interpretation #2

More details

Q: how exactly is dim. reduction done?

A: set the smallest singular values to zero:

1	1	1	0	0
2	2	2	0	0
1	1	1	0	0
5	5	5	0	0
0	0	0	2	2
0	0	0	3	3
0	0	0	1	1

~

1	1	1	0	0
2	2	2	0	0
1	1	1	0	0
5	5	5	0	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

# SVD - Interpretation #3

finds non-zero 'blobs' in a data matrix

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix} \times \begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix} \times \begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$$

# SVD - Interpretation #3

finds non-zero 'blobs' in a data matrix

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ \hline 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix} \times \begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix} \times \begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$$

# SVD - Interpretation #3

- finds non-zero 'blobs' in a data matrix =
- 'communities' (bi-partite cores, here)

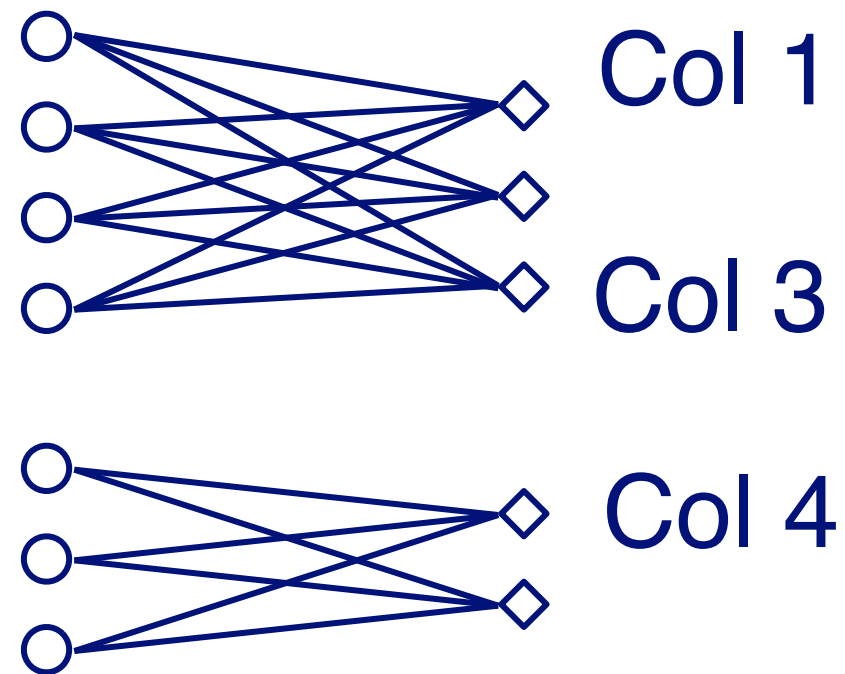
1	1	1	0	0
2	2	2	0	0
1	1	1	0	0
5	5	5	0	0
			2	2
0	0	0	3	3
0	0	0	1	1

Row 1

Row 4

Row 5

Row 7



# SVD - Complexity

$O(n*m*m)$  or  $O(n*n*m)$  (whichever is less)

Faster version, if just want singular values  
or if we want first  $k$  singular vectors  
or if the matrix is sparse [Berry]

No need to write your own!

Available in most linear algebra packages  
(LINPACK, matlab, Splus/R,  
mathematica ...)

Case Study

**How to do queries with LSI?**

# Case Study

## How to do queries with LSI?

For example, how to find documents with 'data'?

Diagram illustrating the LSI query process for finding documents related to 'data'.

**Input Matrix (CS docs vs MD docs):**

	data	info	retrieval	brain	lung
CS docs 1	1	1	1	0	0
CS docs 2	2	2	2	0	0
CS docs 3	1	1	1	0	0
CS docs 4	5	5	5	0	0
MD docs 1	0	0	0	2	2
MD docs 2	0	0	0	3	3
MD docs 3	0	0	0	1	1

**Query Vector (data):**

0.18	0
0.36	0
0.18	0
0.90	0
0	0.53
0	0.80
0	0.27

**Intermediate Matrix:**

9.64	0
0	5.29

**Result Matrix:**

0.58	0.58	0.58	0	0
0	0	0	0.71	0.71

The diagram shows the calculation of the query vector for 'data' and its subsequent multiplication with the intermediate matrix to find related documents.



## Case Study

# How to do queries with LSI?

For example, how to find documents with 'data'?

A: map query vectors into 'concept space' – how?

The diagram illustrates the process of mapping query vectors into a concept space for LSI. It shows a matrix of CS docs (Concept Space documents) and a matrix of MD docs (Medical Document documents) being multiplied by three intermediate matrices to produce a final result.

**CS docs matrix (Green text):**

	data	info	retrieval	brain	lung
1	1	1	1	0	0
2	2	2	2	0	0
1	1	1	1	0	0
5	5	5	5	0	0

**MD docs matrix (Orange text):**

0	0	0	2	2
0	0	0	3	3
0	0	0	1	1

**Intermediate matrices:**

**Matrix 1 (Green text):**

0.18	0
0.36	0
0.18	0
0.90	0
0	0.53
0	0.80
0	0.27

**Matrix 2 (Green text):**

9.64	0
0	5.29

**Matrix 3 (Green text):**

0.58	0.58	0.58	0	0
0	0	0	0.71	0.71

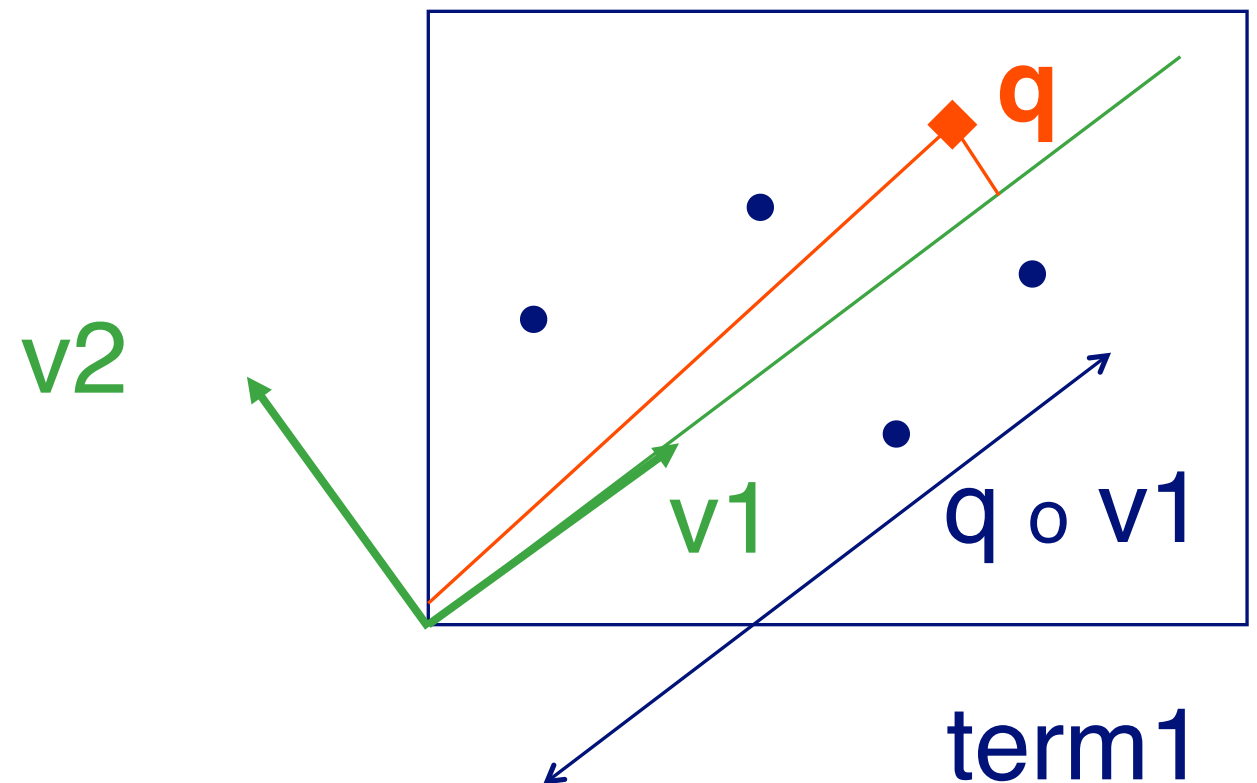
The final result is the product of these three matrices.

## Case Study

# How to do queries with LSI?

For example, how to find documents with ‘data’?  
A: map query vectors into ‘concept space’, using **inner product** (cosine similarity) with each ‘concept’ vector  $v_i$

$$\mathbf{q} = \begin{array}{c} \text{data} \\ \text{info} \\ \text{retrieval} \\ \text{brain} \\ \text{lung} \end{array} \begin{array}{|c|c|c|c|c|} \hline 1 & 0 & 0 & 0 & 0 \\ \hline \end{array}$$

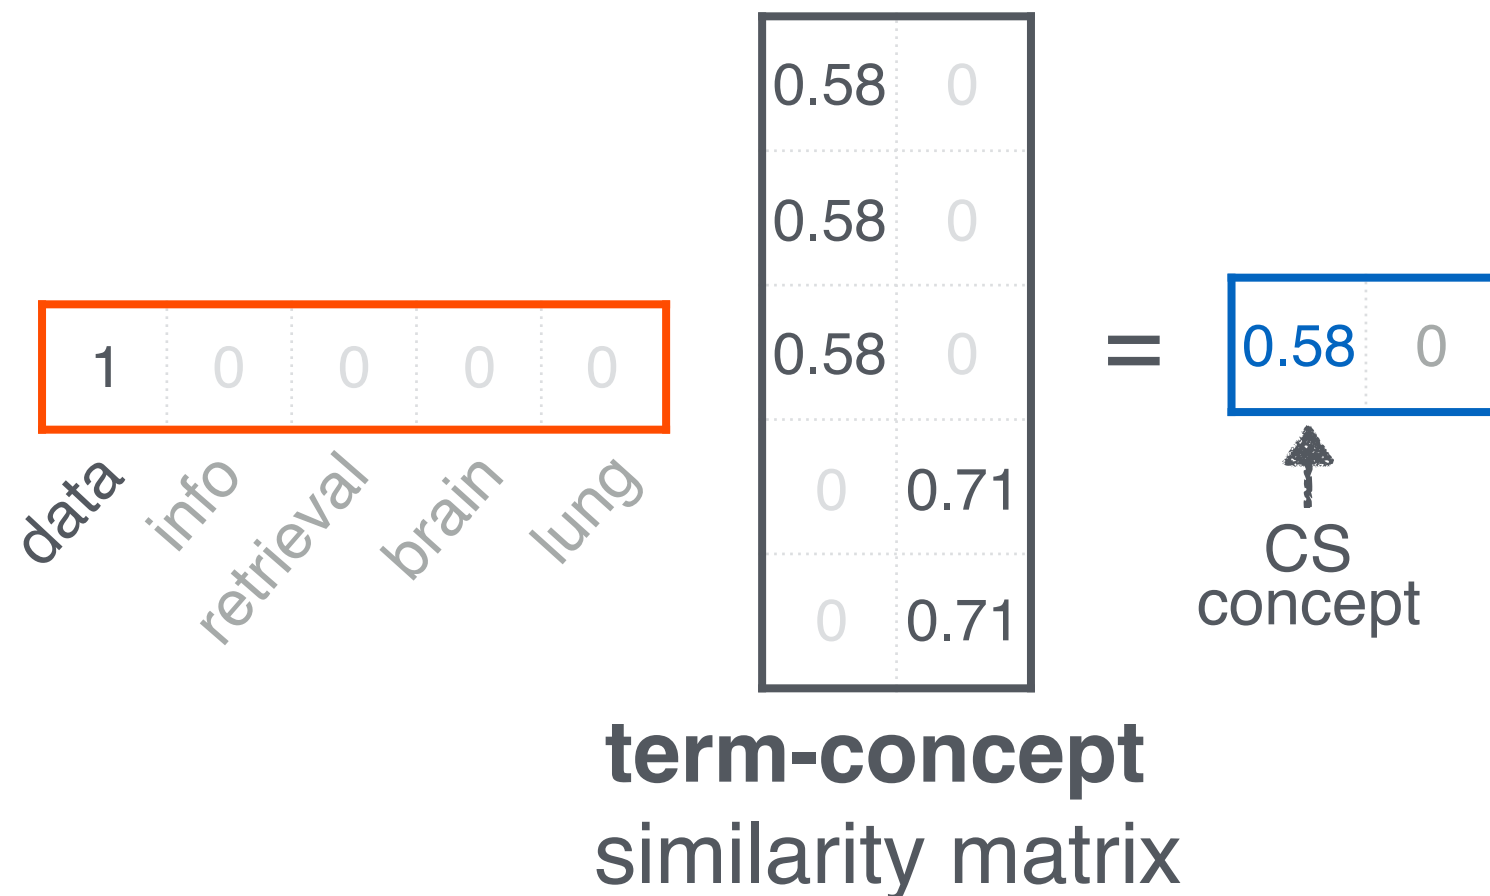


# Case Study

## How to do queries with LSI?

Compactly, we have:

$$\mathbf{q} \mathbf{V} = \mathbf{q}_{\text{concept}}$$



Case Study

**How would the document  
(‘information’, ‘retrieval’) be handled?**

## Case Study

**How would the document  
(‘information’, ‘retrieval’) be handled?**

**SAME!**

$$\mathbf{d} \mathbf{V} = \mathbf{d}_{\text{concept}}$$

0	1	1	0	0
data	info	retrieval	brain	lung

0.58	0
0.58	0
0.58	0
0	0.71
0	0.71

1.16	0
------	---

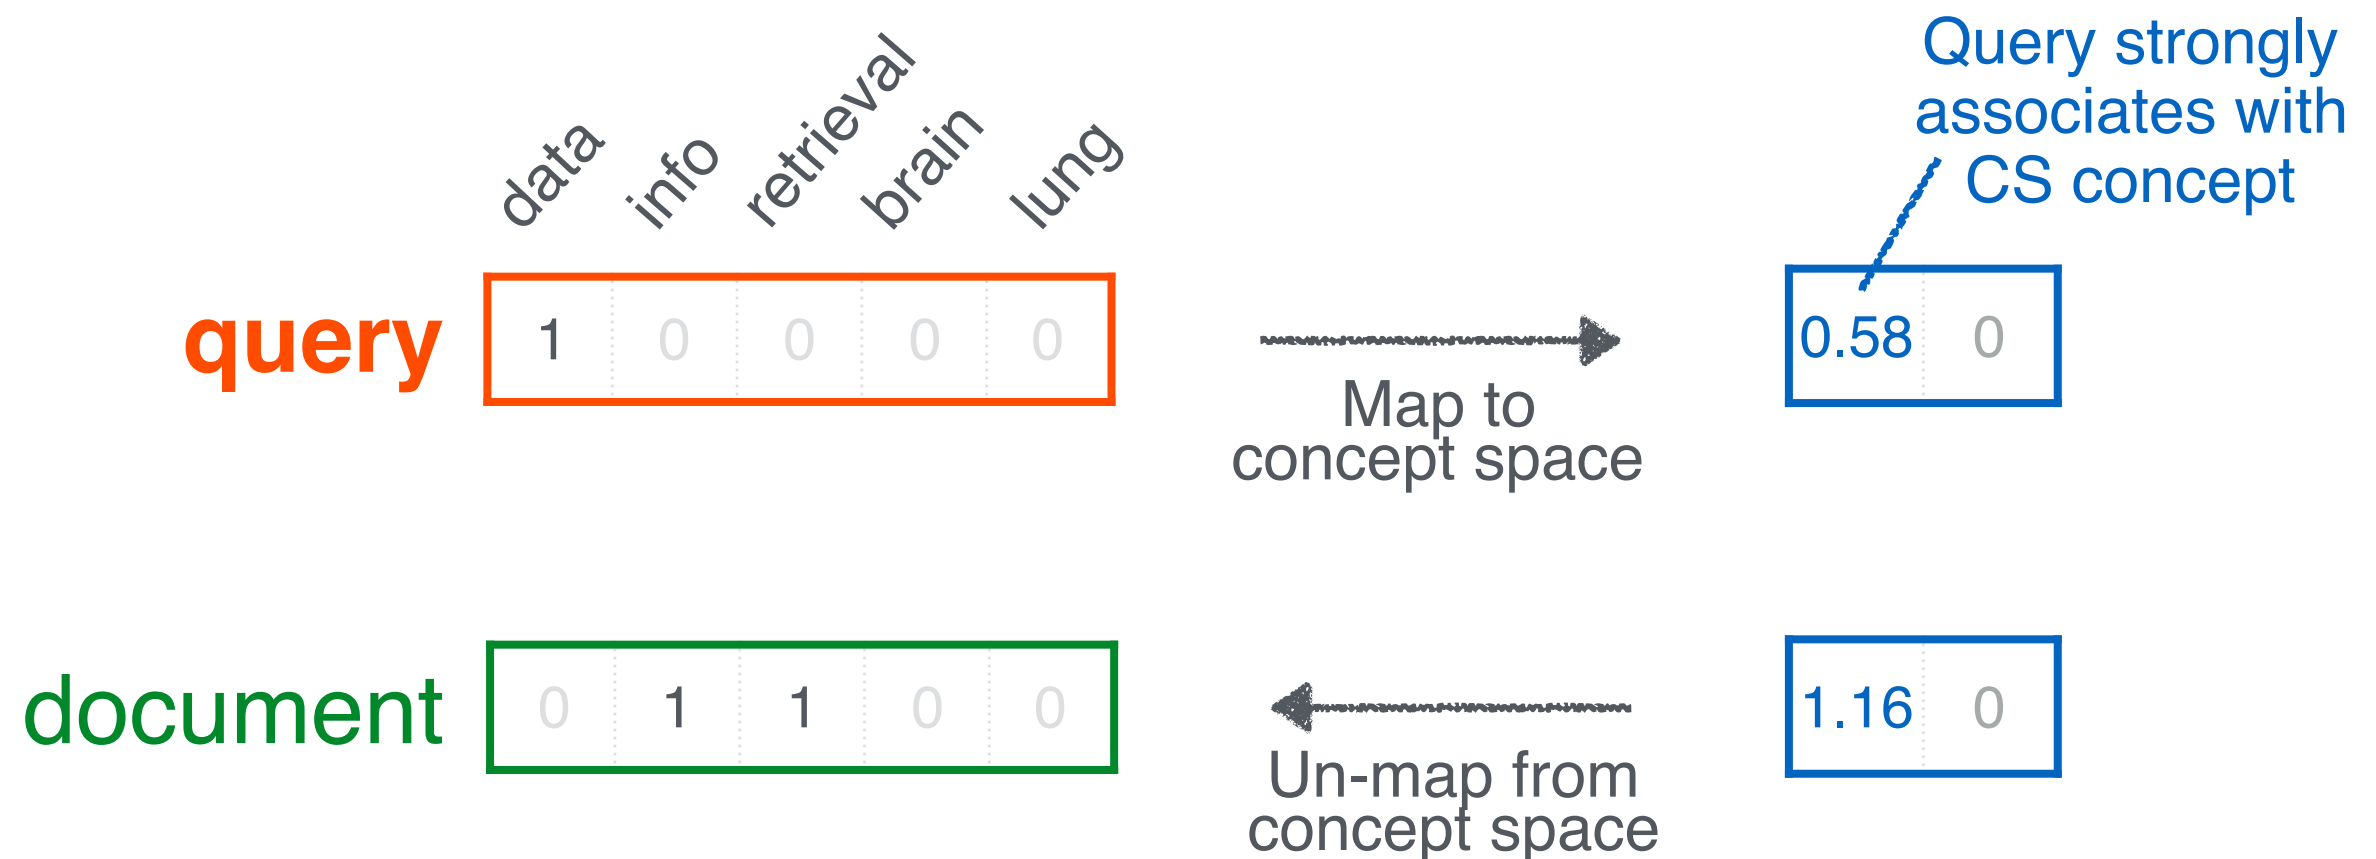
↑  
CS  
concept

**term-concept  
similarity matrix**

# Case Study

## Observation

**Document** ('information', 'retrieval') will be retrieved by **query** ('data'), even though it does not contain 'data'!!



# Case study - LSI

Q1: How to do queries with LSI?

➡ Q2: multi-lingual IR (english query, on spanish text?)

# Case study - LSI

- Problem:
  - given many documents, translated to both languages (e.g., English and Spanish)
  - answer queries across languages



# Case study - LSI

- Solution:  $\sim$  LSI

Diagram illustrating a neural network architecture for a classification task. The input layer consists of 10 nodes, with the first 5 labeled 'data', 'inf', 'retrieval', 'brain', and 'lung', and the next 5 labeled 'informacion', 'datos', and three unlabeled nodes. The hidden layer has 5 nodes. The output layer has 2 nodes, labeled 'CS' and 'MD'. Arrows indicate the flow of information from input to hidden to output. A dashed vertical line separates the input nodes from the hidden nodes.

	data	inf	retrieval	brain	lung	informacion	datos			
CS	1	1	1	0	0	1	1	1	0	0
	2	2	2	0	0	1	2	2	0	0
	1	1	1	0	0	1	1	1	0	0
	5	5	5	0	0	5	5	4	0	0
MD	0	0	0	2	2	0	0	0	2	2
	0	0	0	3	3	0	0	0	2	3
	0	0	0	1	1	0	0	0	1	1

# Switch Gear to **Text Visualization**

# Word/Tag Cloud (still popular?)




<http://www.wordle.net>

**Twitter** | [Tweeps](#) | [Wikipedia](#) | [Custom](#)

Keyword: **cloud**

Go!

# Word Counts (words as bubbles)





# Word Tree





# Phrase Net

Visualize pairs of words satisfying a pattern (“X and Y”)

Select a phrase

word1 and word2

word1 's word2

word1 of the word2

word1 the word2

word1 a word2

word1 at word2

word1 is word2

word1 [space] word2

or enter your own  
\* and \* Submit

Filters

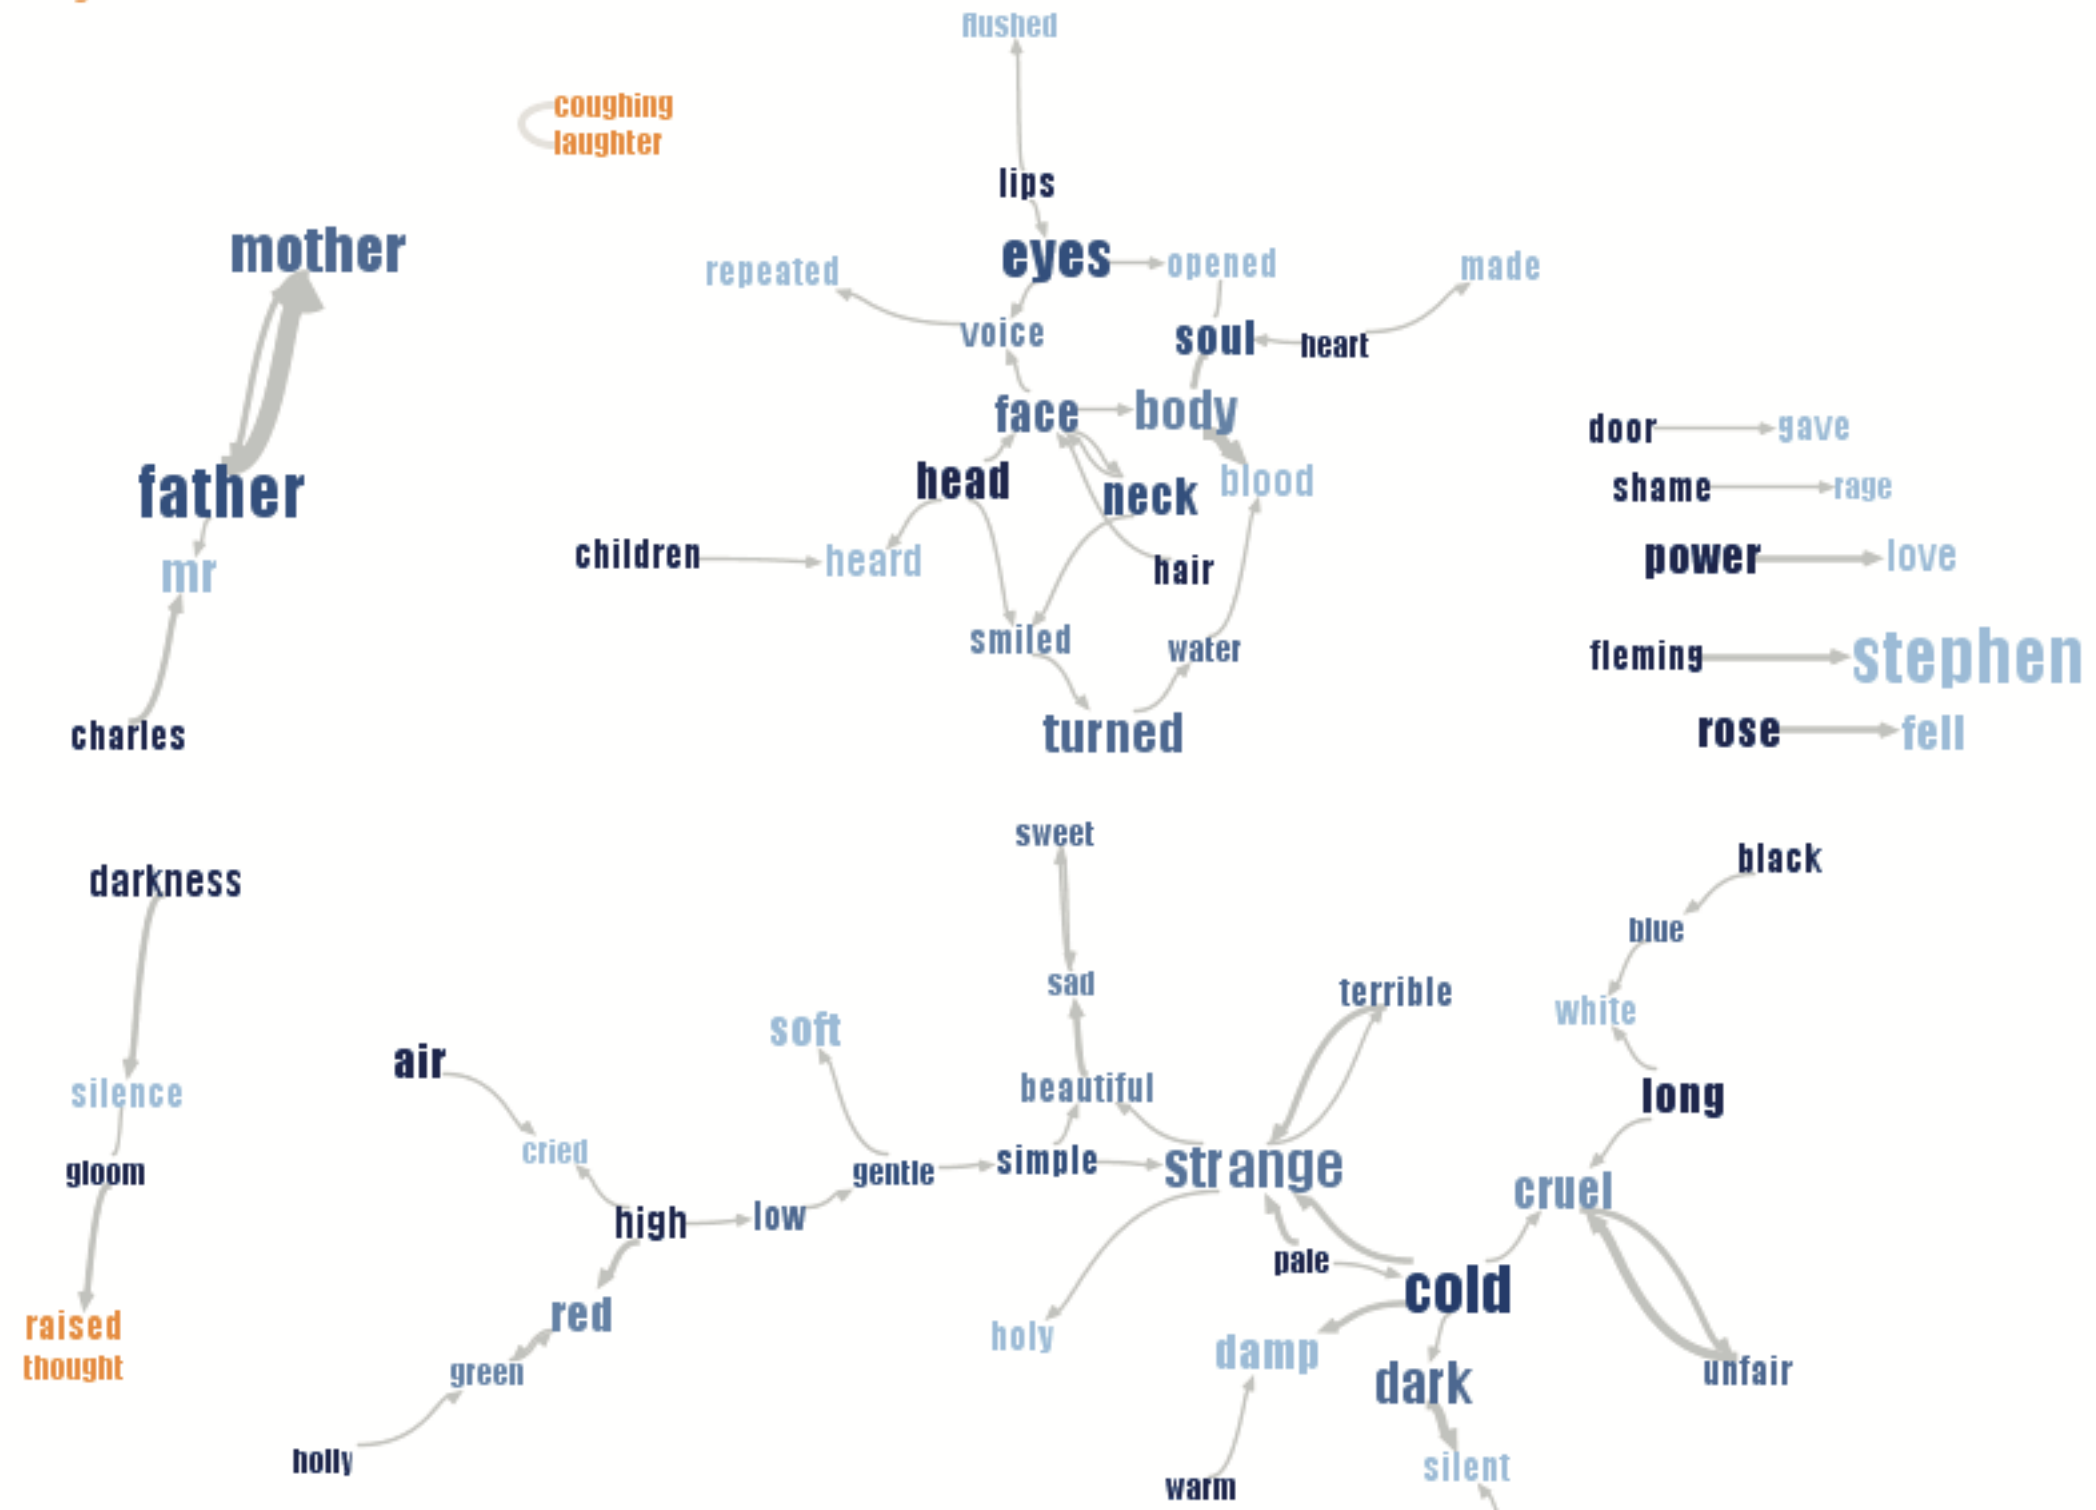
Show top: 100  
Hide common words ☒

Zoom

In Out Reset

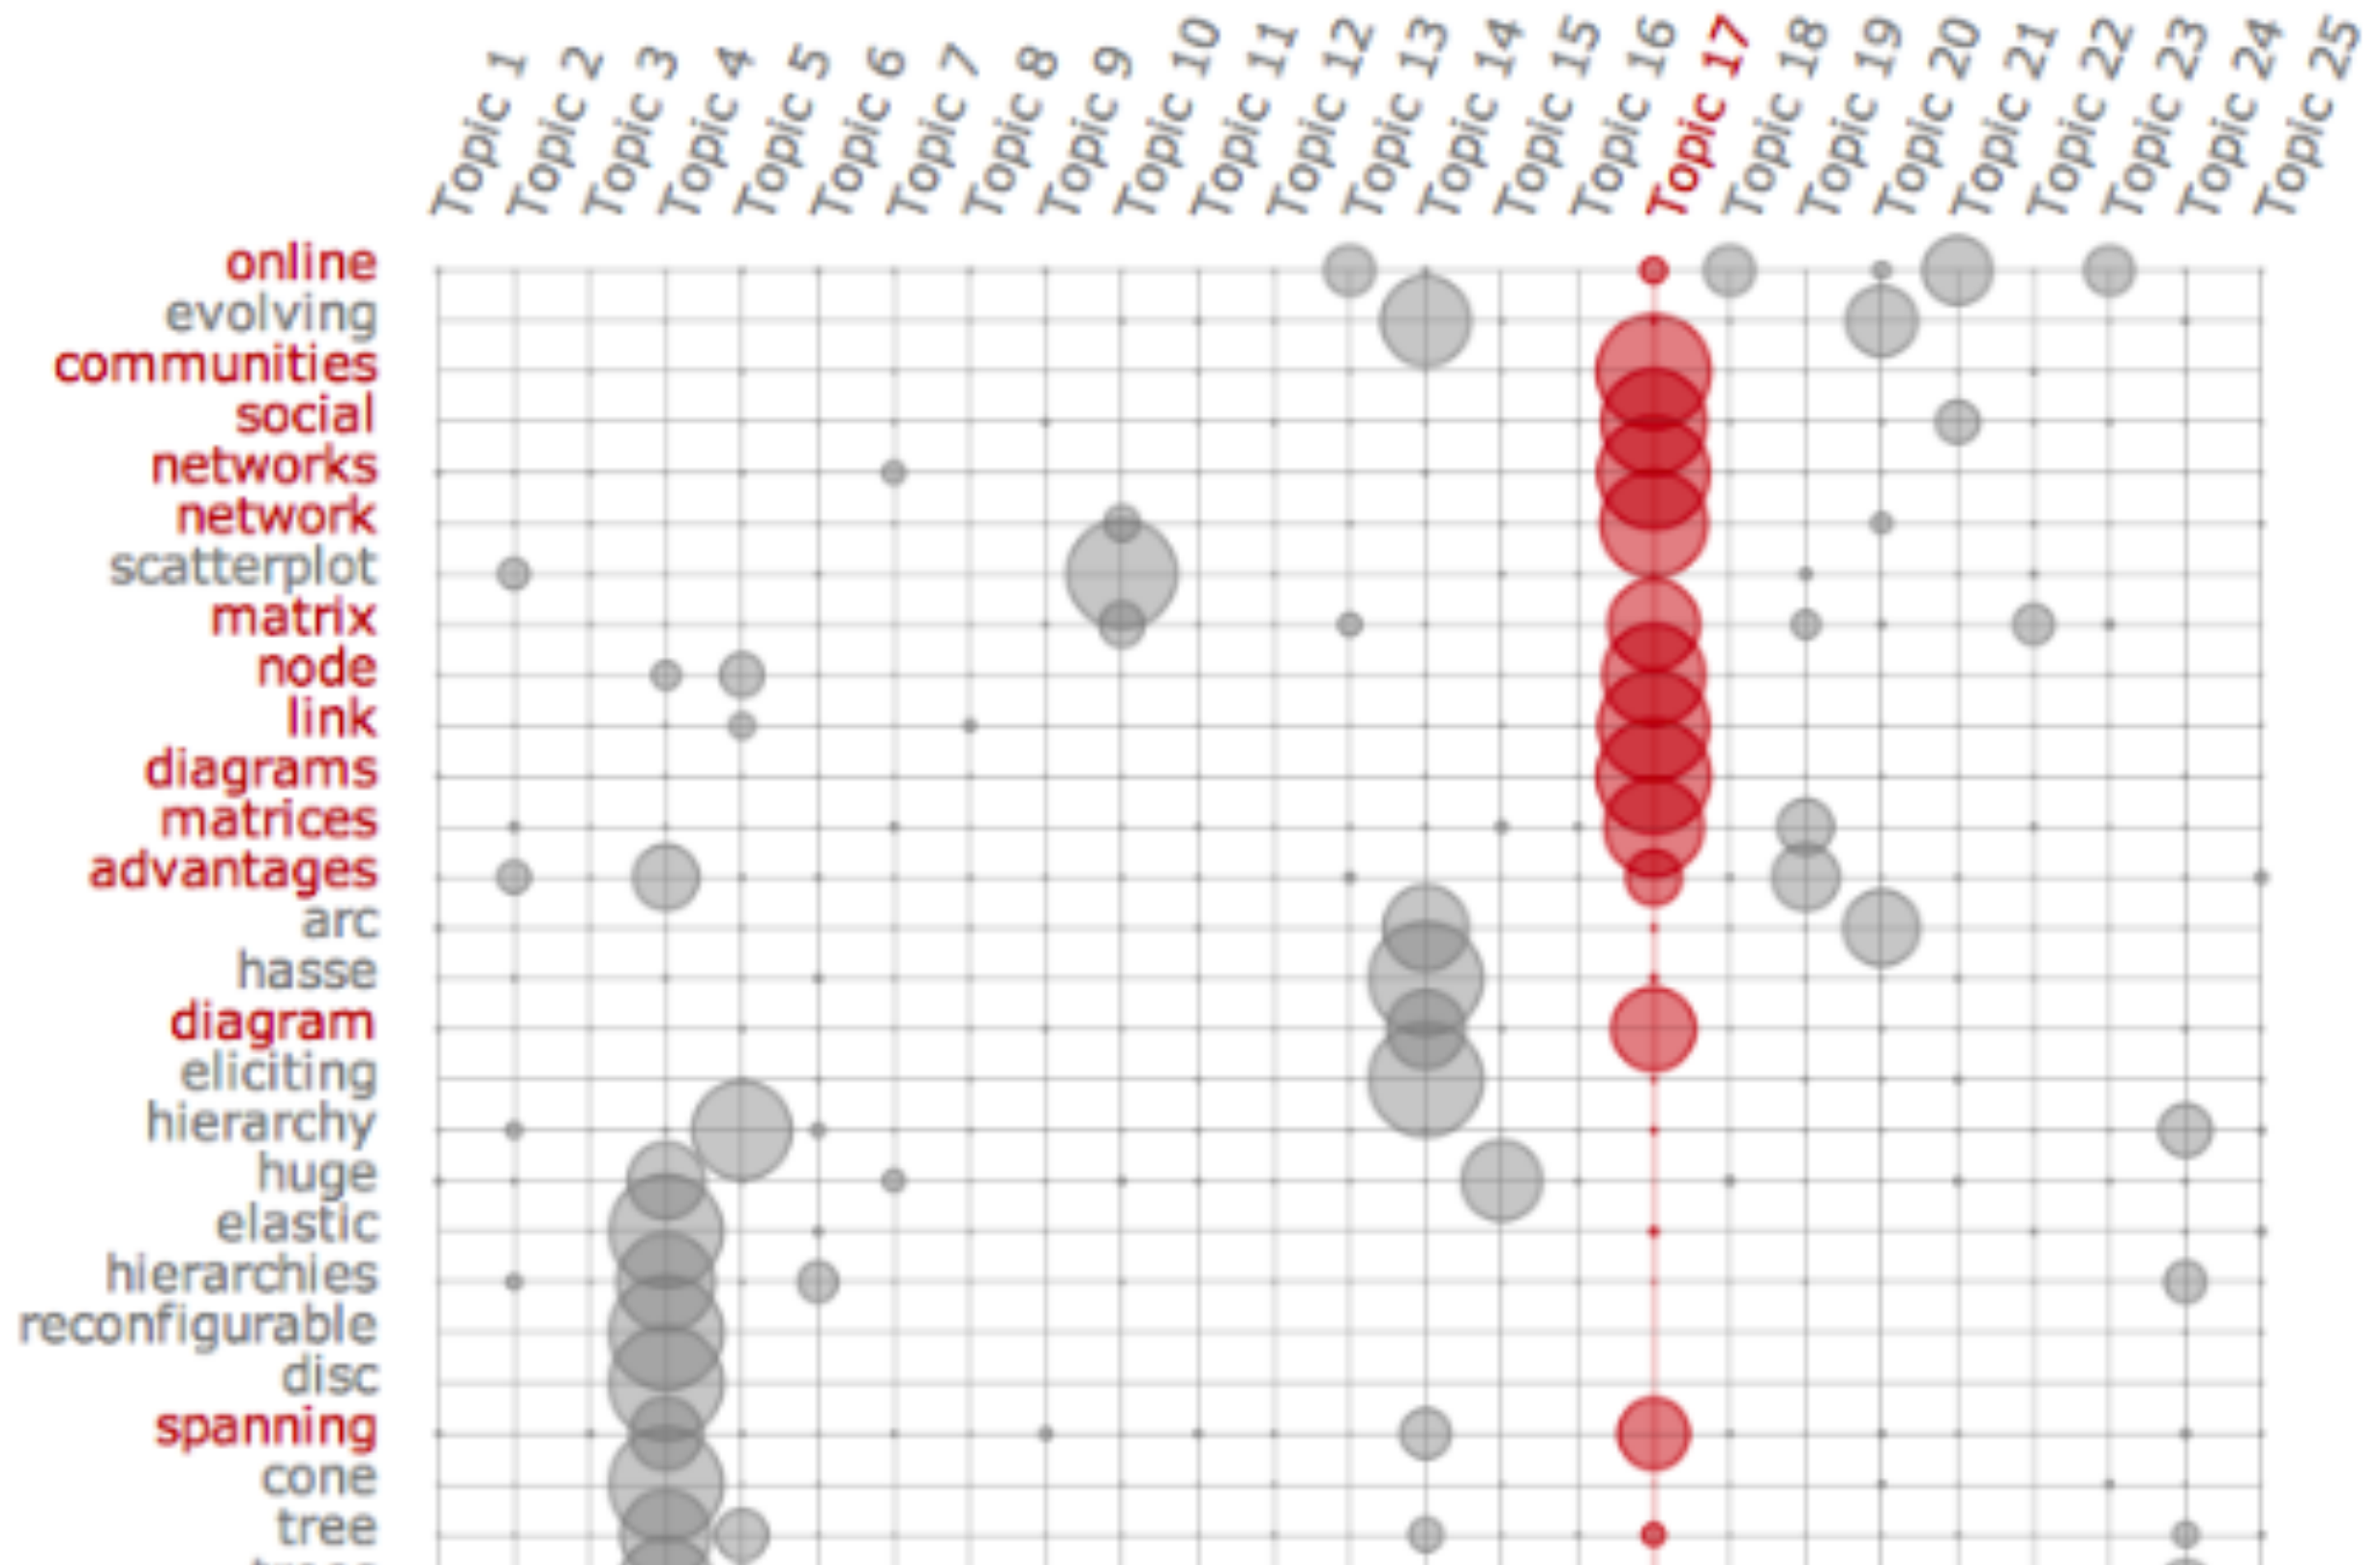
<http://hint.fm/projects/phrasenet/>

Showing 73 of 1719 terms



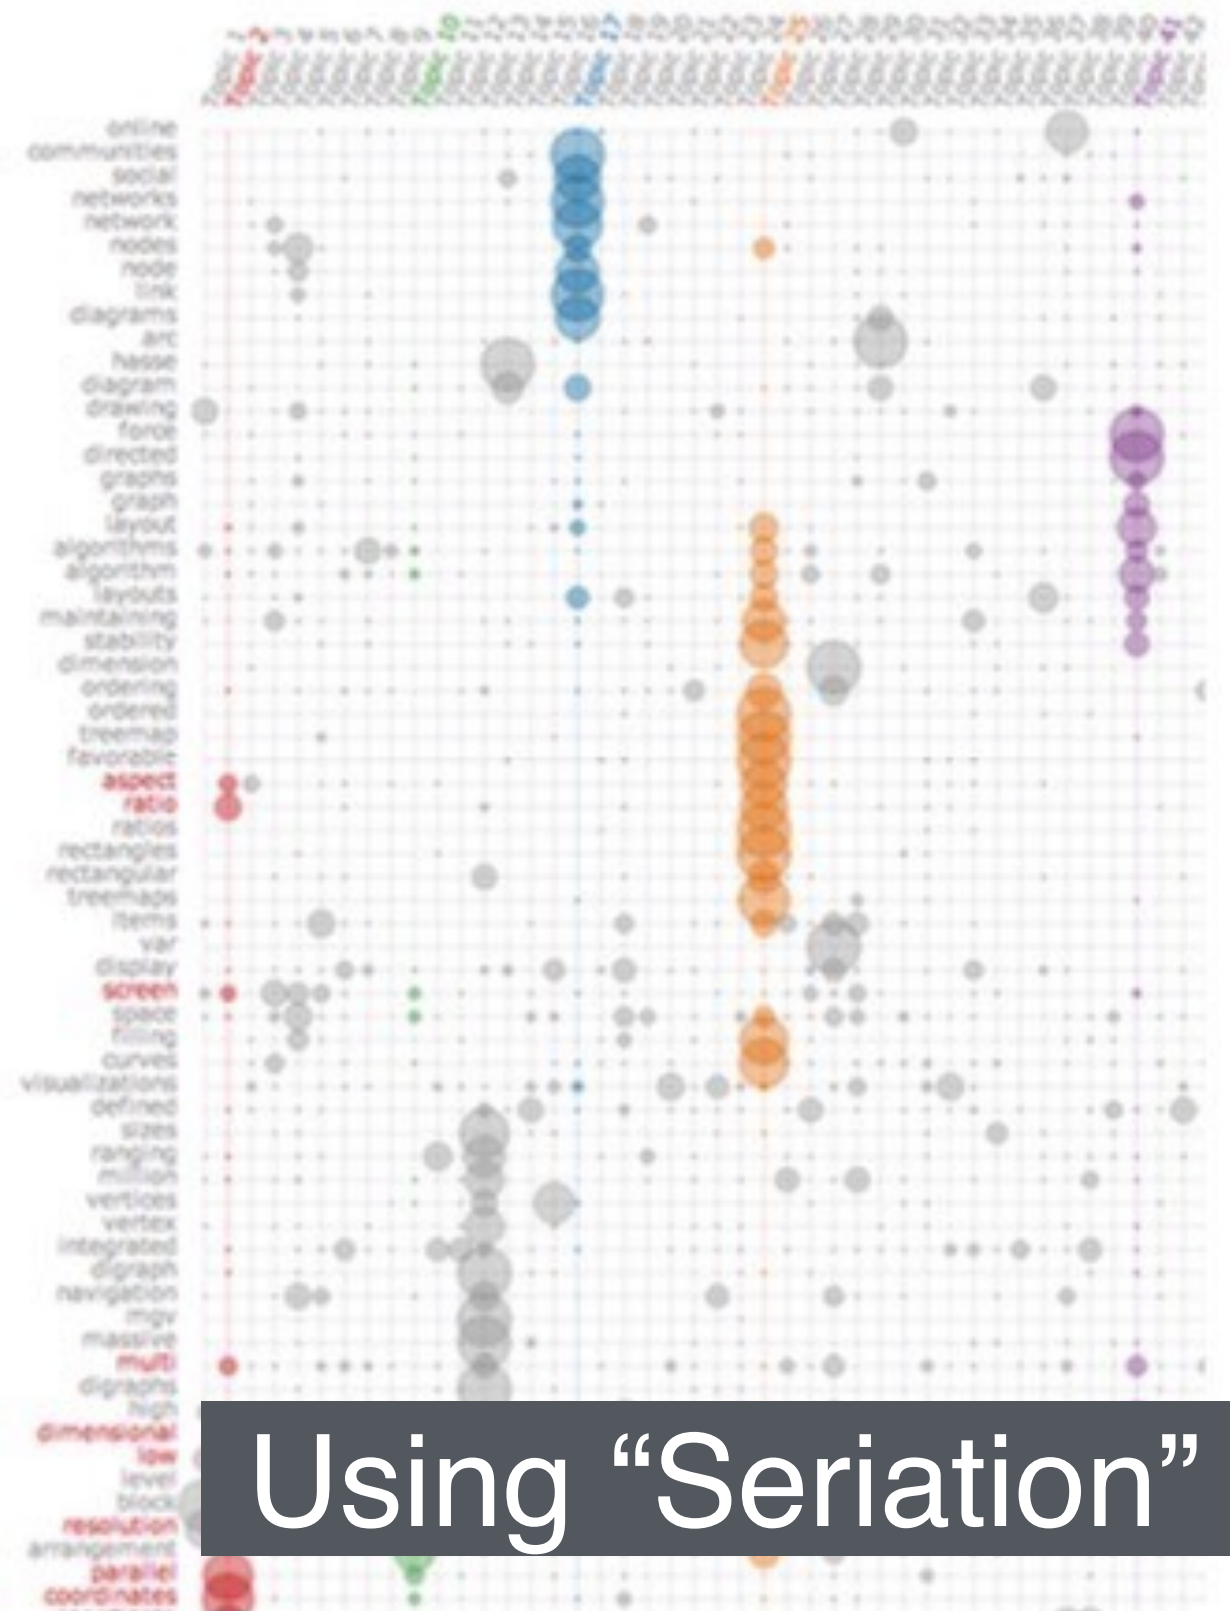

# Termite: Topic Model Visualization

<http://vis.stanford.edu/papers/termite>



# Termite: Topic Model Visualization

<http://vis.stanford.edu/papers/termite>



Using “Seriation”