CSE6242 / CX4242: Data & Visual Analytics

Graphs / Networks

Centrality measures, algorithms, interactive applications

Duen Horng (Polo) Chau

Assistant Professor Associate Director, MS Analytics Georgia Tech

Partly based on materials by Professors Guy Lebanon, Jeffrey Heer, John Stasko, Christos Faloutsos, Parishit Ram (GT PhD alum; SkyTree), Alex Gray

Centrality

= "Importance"

Why Node Centrality?

What can we do if we can rank all the nodes in a graph (e.g., Facebook, LinkedIn, Twitter)?

Why Node Centrality?

What can we do if we can rank all the nodes in a graph (e.g., Facebook, LinkedIn, Twitter)?

- Find celebrities or influential people in a social network (Twitter)
- Find "gatekeepers" who connect communities (headhunters love to find them on LinkedIn)
- What else?







Why Node Centrality?

Helps graph analysis, visualization, understanding, e.g.,

- Let us rank nodes, group or study them by centrality
- Only show subgraph formed by the top 100 nodes, out of the millions in the full graph
 - Similar to google search results (ranked, and they only show you 10 per page)
- Most graph analysis packages already have centrality algorithms implemented. Use them!

Can also compute edge centrality. Here we focus on node centrality.

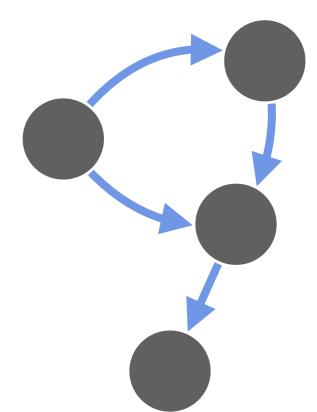
Degree Centrality (easiest)

Degree = number of neighbors

- For directed graphs
 - In degree = No. of incoming edges
 - Out degree = No. of outgoing edges



- Algorithms?
 - Sequential scan through edge list
 - What about for a graph stored in SQLite?



1, 2

1, 3

2, 4

3, 2

Computing Degrees using SQL

Recall simplest way to store a graph in SQLite:

```
edges(source_id, target_id) 1,3
```

- 1. If slow, first create index for each column
- 2. Use group by statement to find out degrees

```
select count(*) from edges group by source_id;
```

2, 4

3, 2

Betweenness Centrality

High betweenness = "gatekeeper"



$$= \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_{st}} - \sum_{s \neq v \neq t \in V} \frac{\sigma_{st}(v)}{\sigma_$$

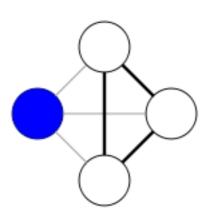
Number of shortest paths between s and t that **goes through v**

Number of shortest paths between s and t

= how often a node serves as the "bridge" that connects two other nodes.

(Local) Clustering Coefficient

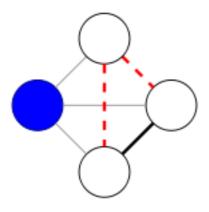
A node's clustering coefficient is a measure of how close the node's neighbors are from forming a clique.



$$\mathsf{c}=\mathsf{1}$$

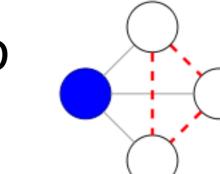
1 = neighbors form a clique





$$c = 1/3$$

(Assuming undirected graph)



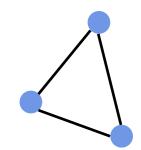
"Local" means it's for a node; can also compute a graph's "global" coefficient

Computing Clustering Coefficients...

Requires triangle counting

Real social networks have a lot of triangles

Friends of friends are friends



Triangles are expensive to compute

(neighborhood intersections; several approx. algos)

Can we do that quickly?

Algorithm details:

Faster Clustering Coefficient Using Vertex Covers http://www.cc.gatech.edu/~ogreen3/_docs/2013VertexCoverClusteringCoefficients.pdf

Super Fast Triangle Counting [Tsourakakis ICDM 2008]



details

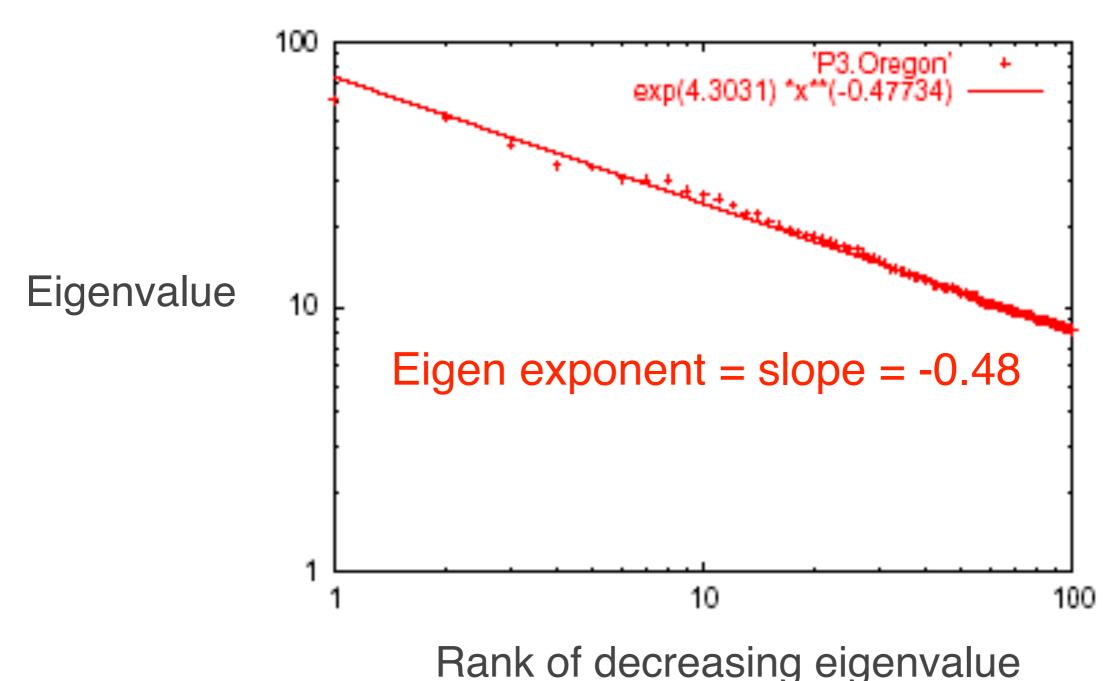
But: triangles are expensive to compute (3-way join; several approx. algos) Q: Can we do that quickly?

A: Yes!

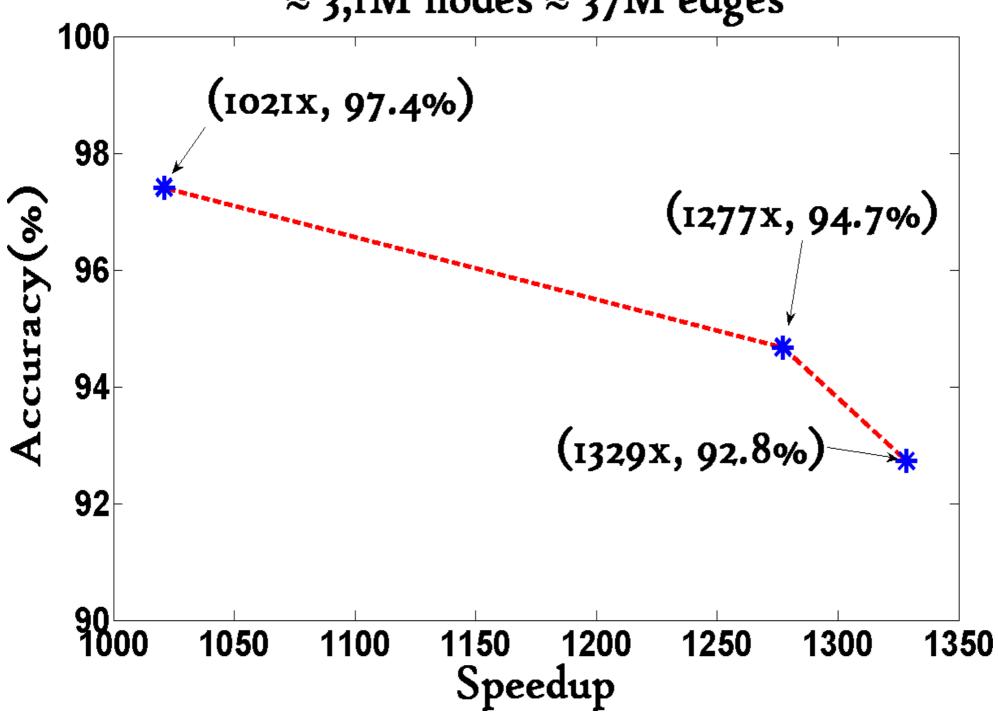
#triangles = 1/6 Sum (λ_i^3)

(and, because of skewness, we only need the top few eigenvalues!

Power Law in Eigenvalues of Adjacency Matrix



Wikipedia graph 2006-Nov-04 ≈ 3,1M nodes ≈ 37M edges



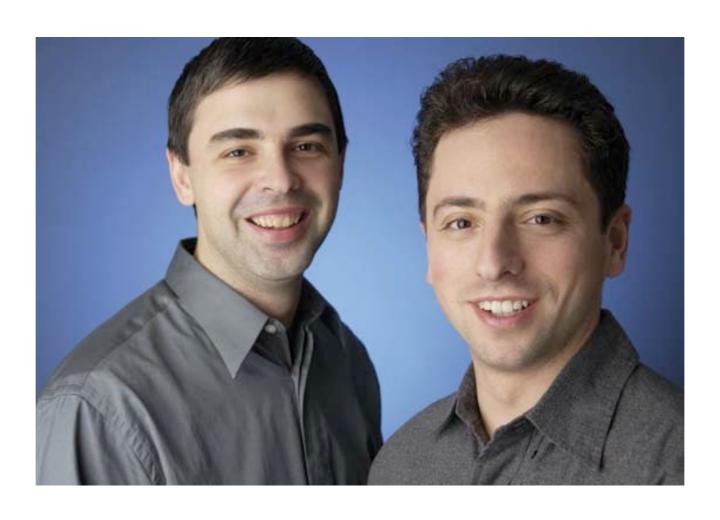
1000x+ speed-up, >90% accuracy

More Centrality Measures...

- Degree
- Betweenness
- Closeness, by computing
 - Shortest paths
 - "Proximity" (usually via random walks) used successfully in a lot of applications
- Eigenvector

• ...

PageRank (Google)



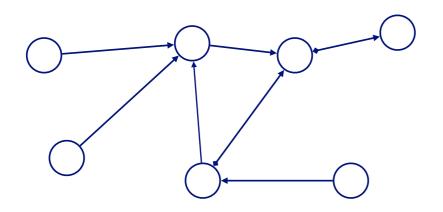
Larry Page

Sergey Brin

Brin, Sergey and Lawrence Page (1998). Anatomy of a Large-Scale Hypertextual Web Search Engine. 7th Intl World Wide Web Conf.

PageRank: Problem

Given a directed graph, find its most interesting/central node



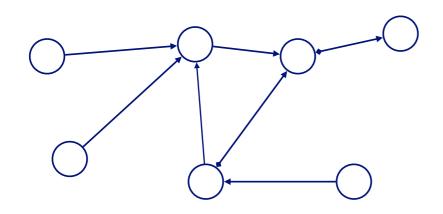
A node is important, if it is connected with important nodes (recursive, but OK!)

PageRank: Solution

Given a directed graph, find its most interesting/central node

Proposed solution:

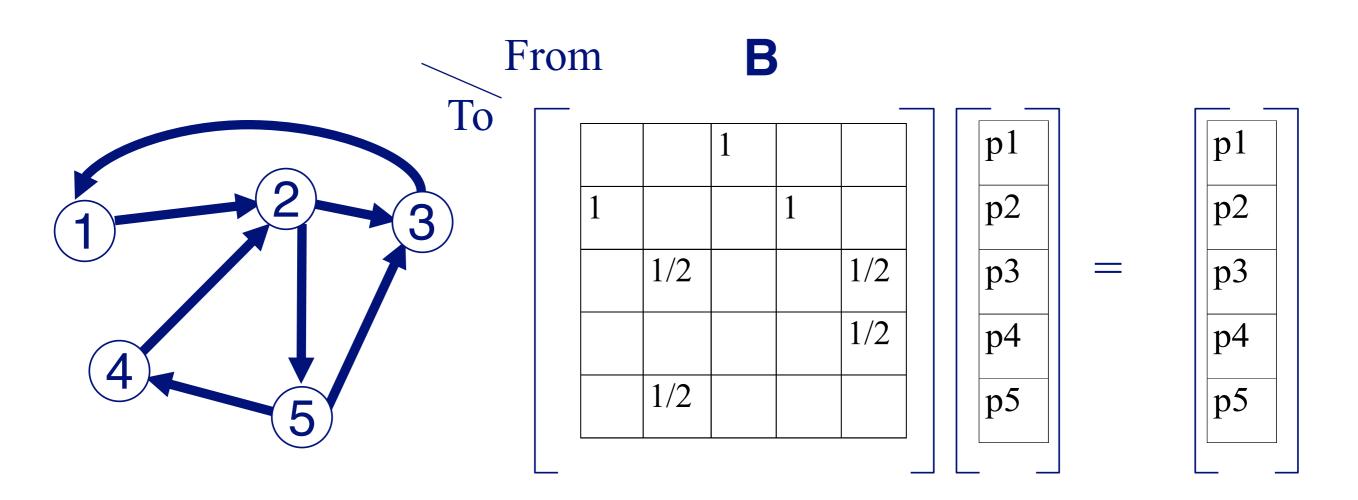
use **random walk**; spot most "popular" node (-> steady state probability (ssp))



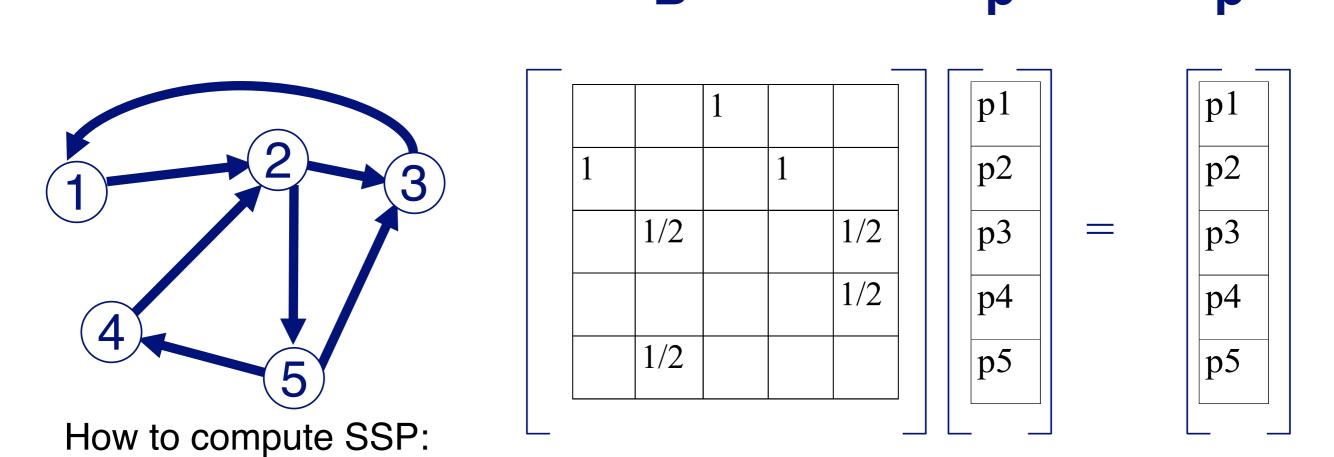
"state" = webpage

A node has high ssp, if it is connected with high ssp nodes (recursive, but OK!)

Let **B** be the transition matrix: transposed, column-normalized



$$Bp = p$$



https://fenix.tecnico.ulisboa.pt/downloadFile/3779579688473/6.3.pdf http://www.sosmath.com/matrix/markov/markov.html

- B p = 1 * p
- Thus, p is the eigenvector that corresponds to the highest eigenvalue (=1, since the matrix is column-normalized)
- Why does such a p exist?
 - p exists if **B** is nxn, nonnegative, irreducible[Perron–Frobenius theorem]

- In short: imagine a particle/person randomly moving along the edges/links
- Compute its steady-state probability (ssp)

Full version of algorithm:

With occasional random jumps to any nodes

Why? To make the matrix irreducible.

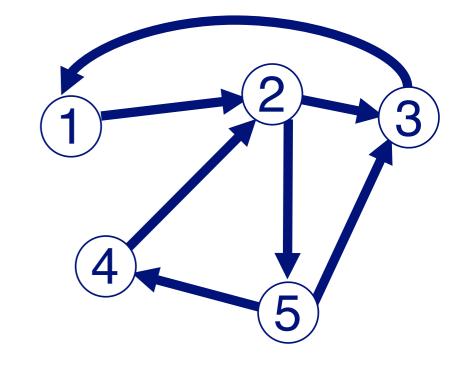
Irreducible = from any state (node), there's non-zero probability to reach any other state (node)

Full Algorithm

With probability 1-c, fly-out to a random node

Then, we have

$$p = c B p + (1-c) 1$$



$$p = (1-c) [I - c B]^{-1} 1$$
1
...
1

How to compute PageRank for huge matrix?

Use the power iteration method

http://en.wikipedia.org/wiki/Power_iteration

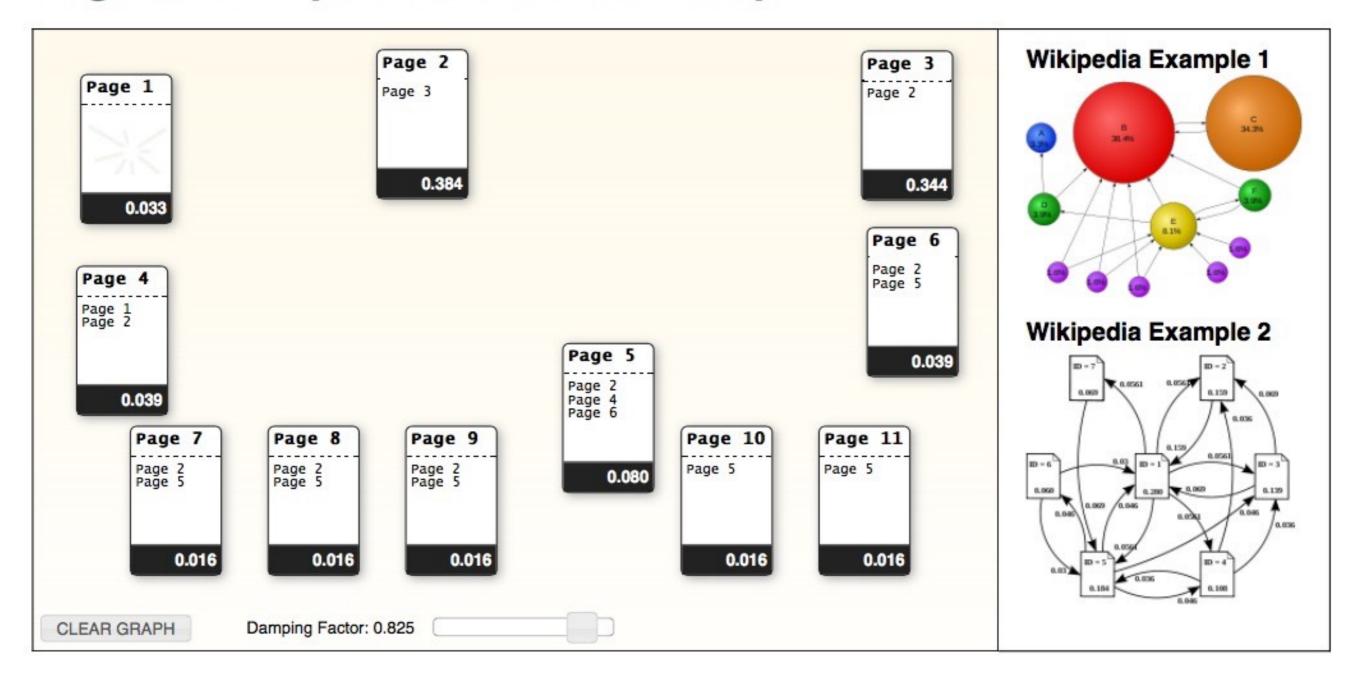
$$p = c B p + (1-c) 1$$

p1	
p2	
p3	
p4	
p5	

				-
		1		
1			1	
	1/2			1/2
				1/2
	1/2			

p1 **p**2 p3 p4 **p5**

PageRank Explained with Javascript



Also great for checking the correctness of your PageRank Implementation.

PageRank for graphs (generally)

You can run PageRank on any graphs

All you need are the graph edges!

Should be in your algorithm "toolbox"

- Better than degree centrality
- Fast to compute for large graphs, runtime linear in the number of edges, O(E)

But can be "misled" (Google Bomb)

How?

Personalized PageRank

Intuition: not all pages are equal, some more relevant to some people

Goal: rank pages in a way that those more relevant to you will be ranked higher

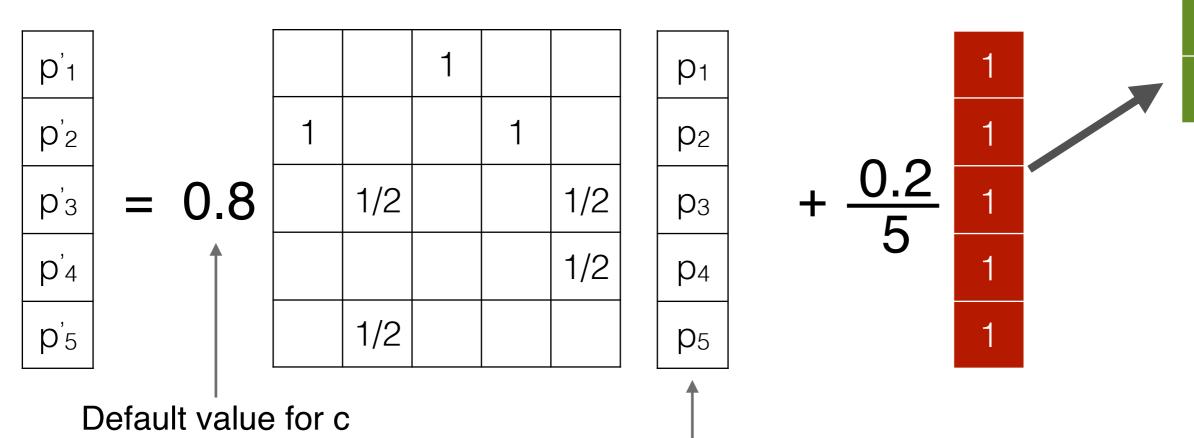
How? Make just one small change to PageRank

Personalized PageRank

With probability 1-c, fly-out to

a random node some preferred nodes

$$p' = c B p + (1-c) 1$$



Can initialize this vector to any non-zero vector, e.g., all "1"s

0

1

0

0

1

Why Learn Personalized PageRank?

For recommendation

- If I like webpage A, what else do I like?
- If I bought product A, what other products would I also buy?

Visualizing and interacting with large graphs

 Instead of visualizing every single nodes, visualize the most important ones

Very flexible — works on any graph

Related "guilt-by-association" / diffusion techniques

- Personalized PageRank
 - (= Random Walk with Restart)
- "Spreading activation" or "degree of interest" in Human-Computer Interaction (HCI)
- Belief Propagation
 (powerful inference algorithm, for fraud detection, image segmentation, error-correcting codes, etc.)

Why are these algorithms popular?

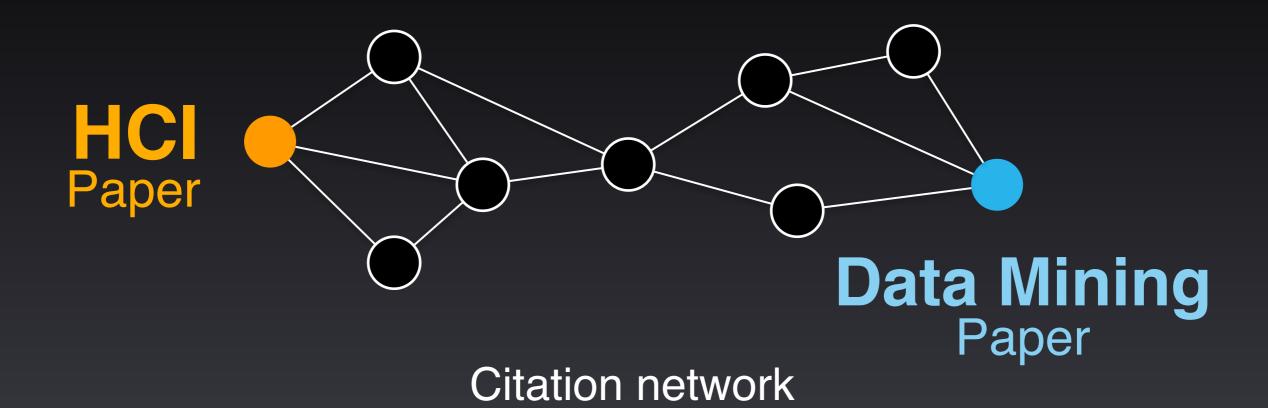
- Intuitive to interpret uses "network effect", homophily
- Easy to implement math is relatively simple (mainly matrixvector multiplication)
- Fast run time linear to #edges, or better
- Probabilistic meaning

Human-In-The-Loop Graph Mining

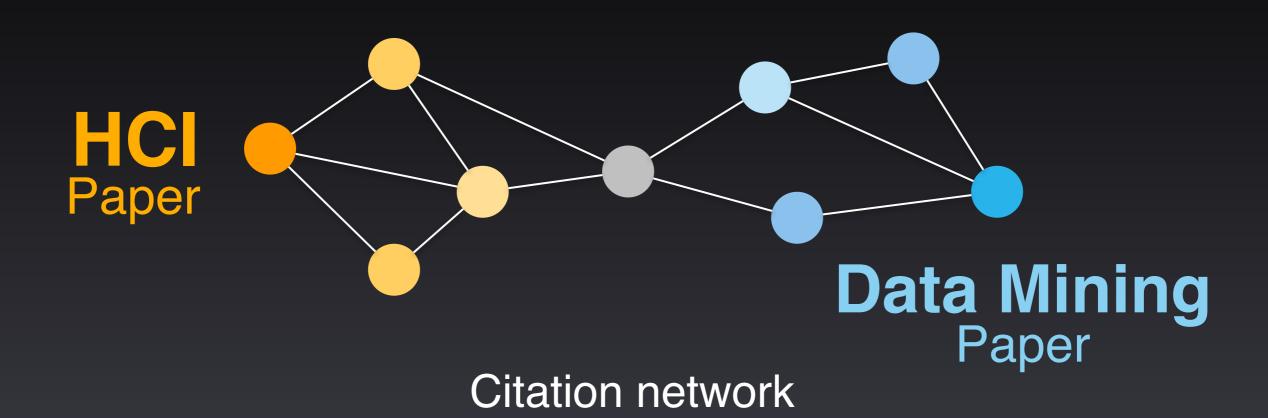
Apolo: Machine Learning + Visualization CHI 2011

Apolo: Making Sense of Large Network Data by Combining Rich User Interaction and Machine Learning

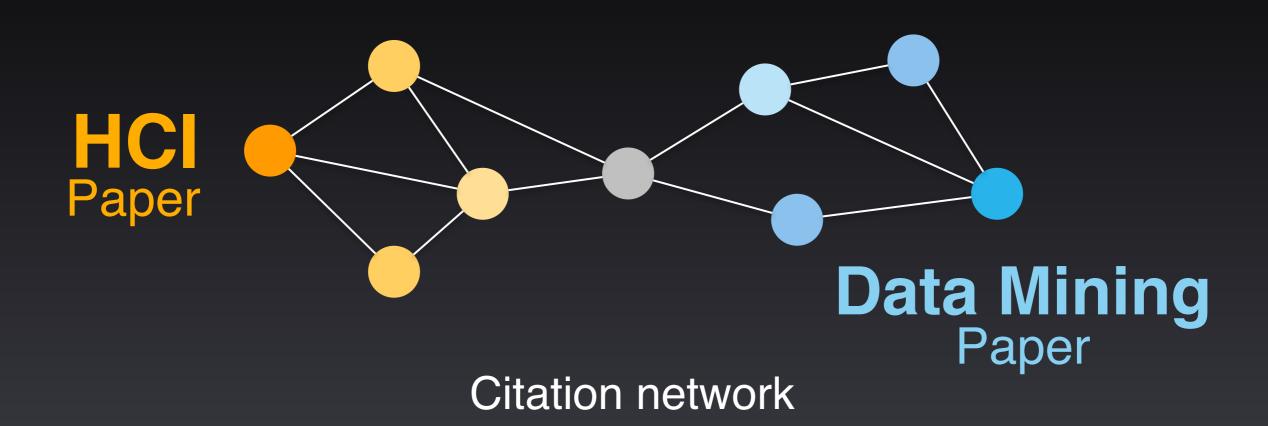
Finding More Relevant Nodes



Finding More Relevant Nodes



Finding More Relevant Nodes

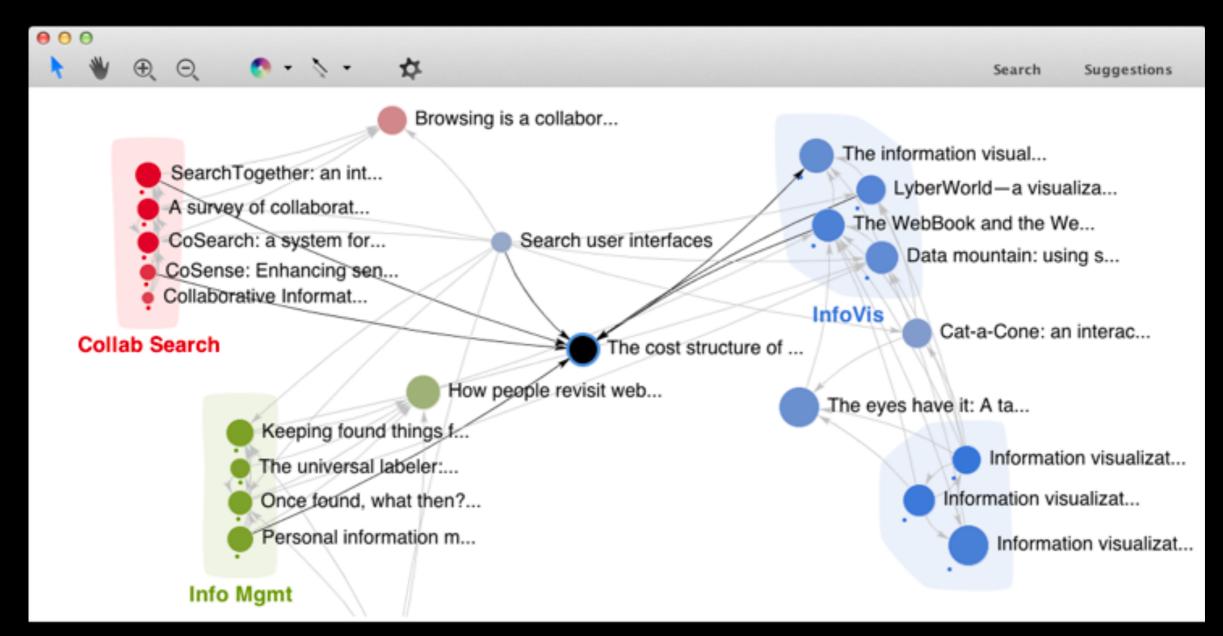


Apolo uses guilt-by-association (Belief Propagation, similar to personalized PageRank)

Demo: Mapping the Sensemaking Literature

Nodes: 80k papers from Google Scholar (node size: #citation)

Edges: 150k citations



The cost structure of sensemaking

Russell, D.M. and Stefik, M.J. and Pirolli, P. and Card, S.K.

245 citations 8 versions

PDF 1993

For The cost structure of sensemaking







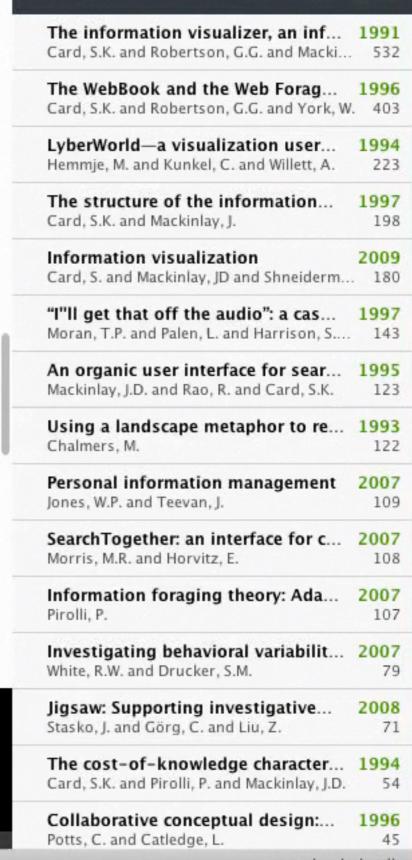


The cost structure of sen...

The cost structure of sensemaking

Russell, D.M. and Stefik, M.J. and Pirolli, P. and Card, S.K.

245 citations 8 versions 1993





For The cost structure of sensemaking



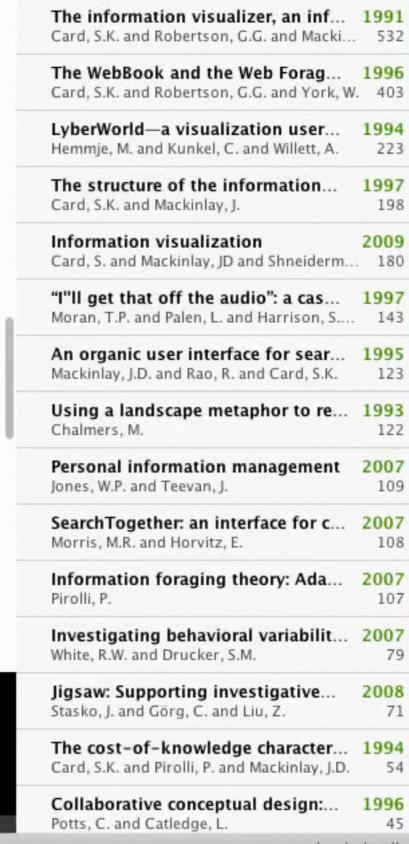


The cost structure of sen...

The cost structure of sensemaking

Russell, D.M. and Stefik, M.J. and Pirolli, P. and Card, S.K.

245 citations 8 versions 1993

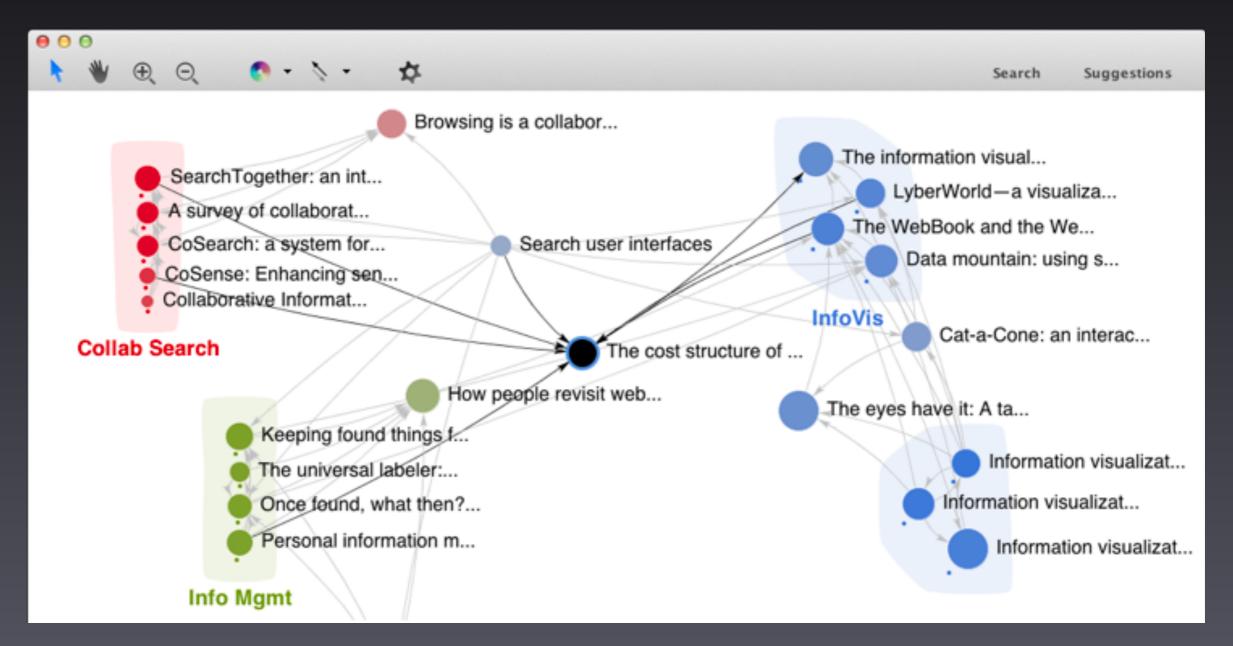


Key Ideas (Recap)



Specify exemplars

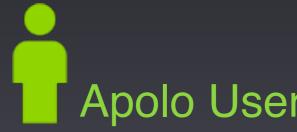
Find other relevant nodes (BP)



Apolo's Contributions

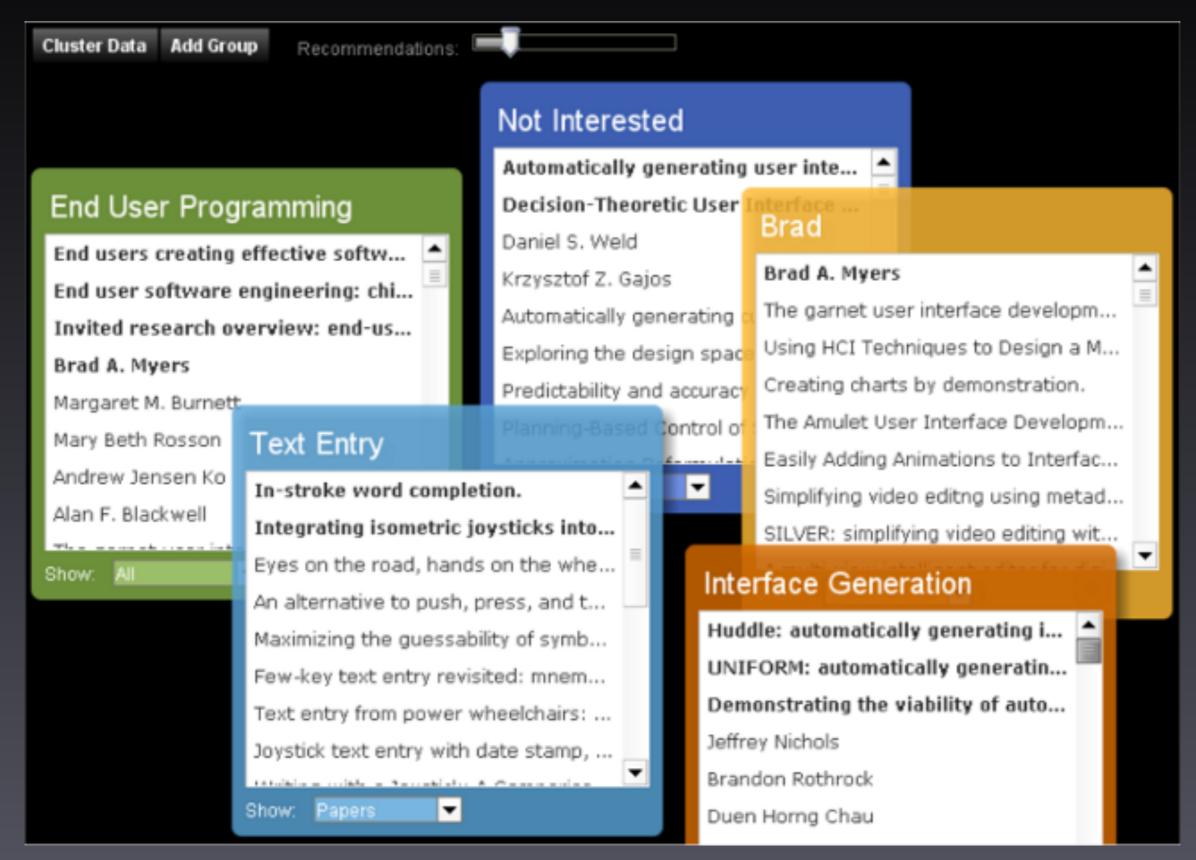
Human + Machine

It was like having a partnership with the machine.

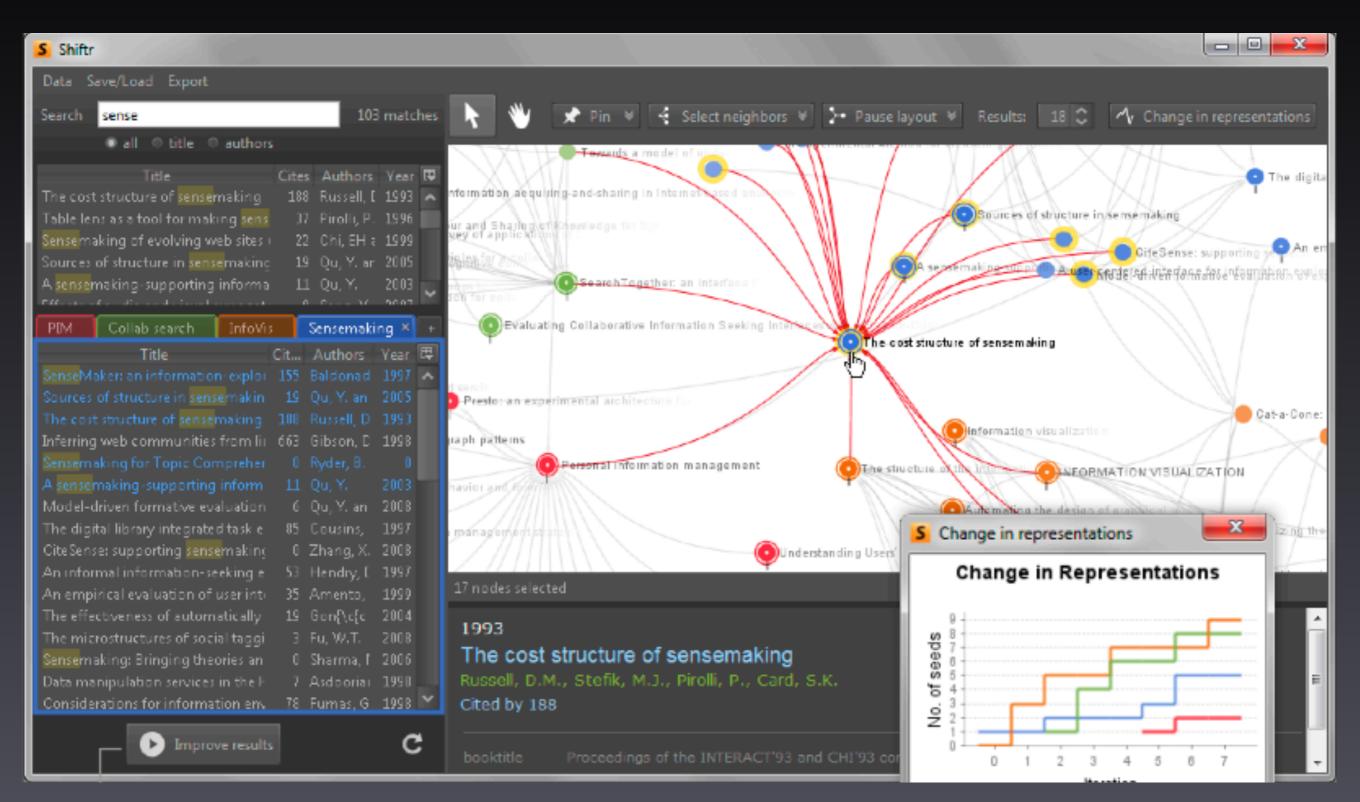


2 Personalized Landscape

Apolo 2009

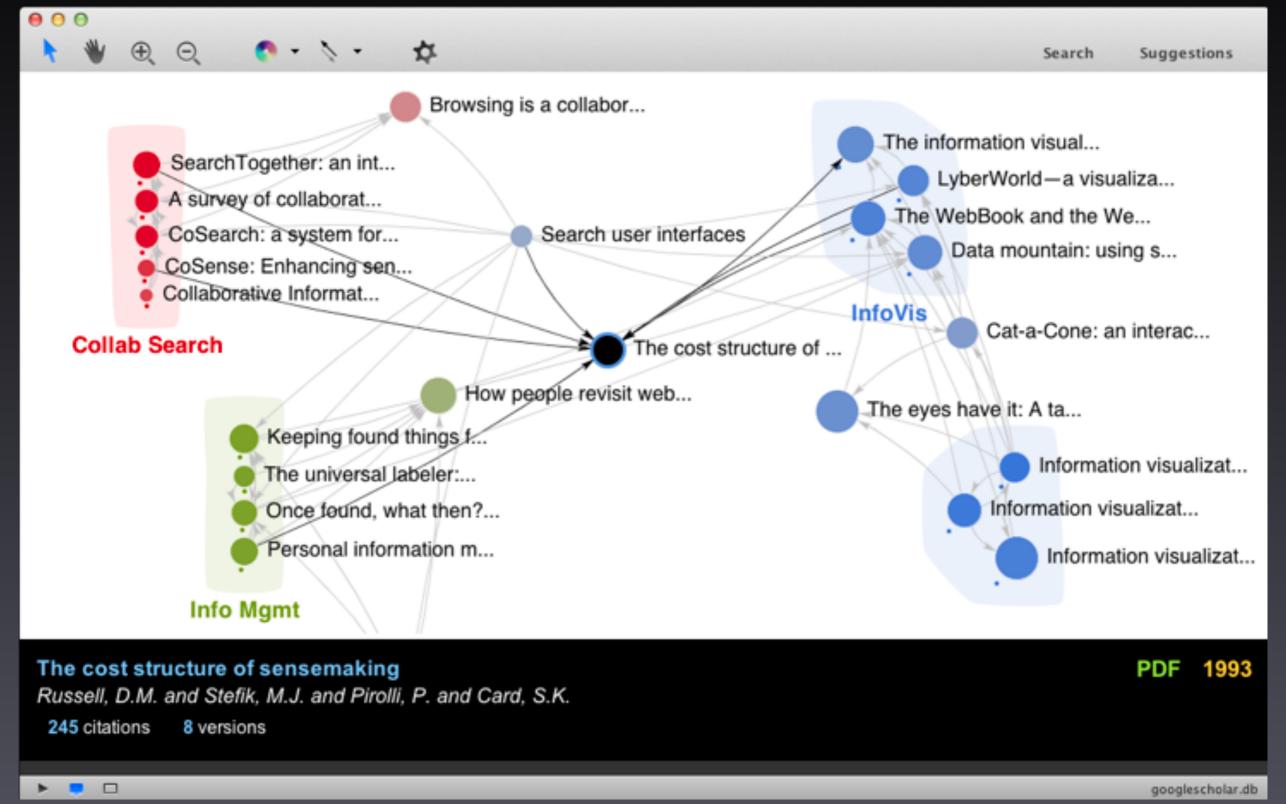


Apolo 2010



Apolo 2011

22,000 lines of code. Java 1.6. Swing. Uses SQLite3 to store graph on disk



User Study

Used citation network

Task: Find related papers for 2 sections in a survey paper on *user interface*

- Model-based generation of UI
- Rapid prototyping tools

Past, Present and Future of User Interface Software Tools

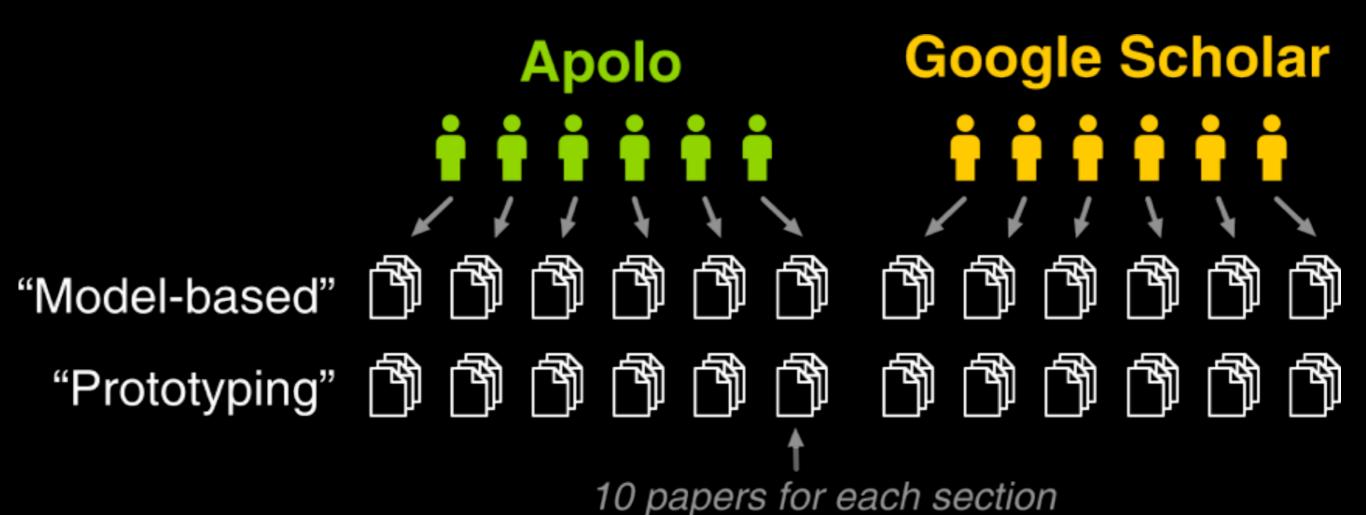
Brad Myers, Scott E. Hudson, and Randy Pausch

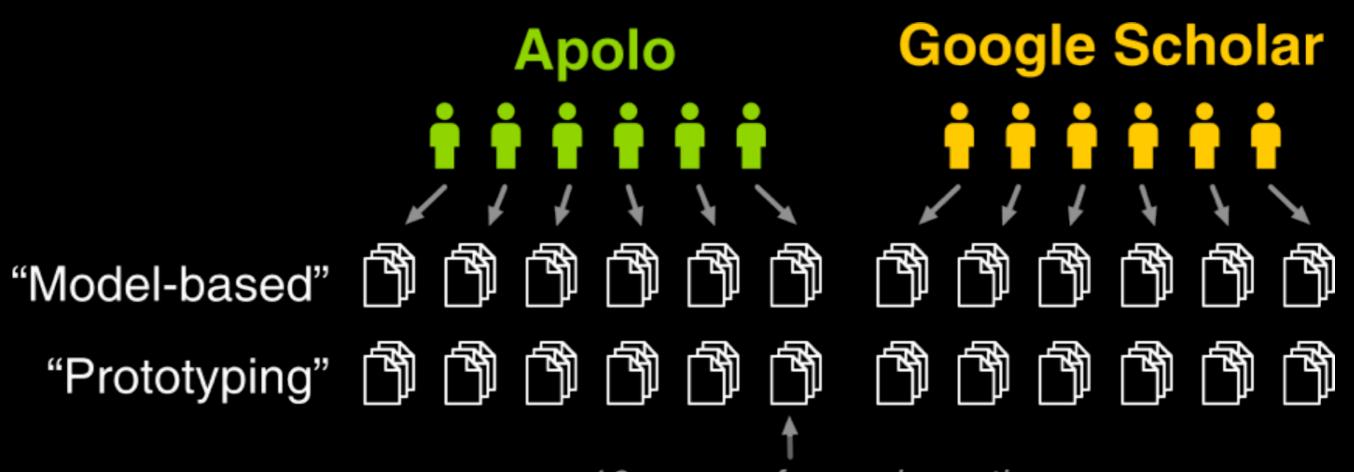
Human Computer Interaction Institute
School of Computer Science
Carnegie Mellon University
Pittelweek PA 15213-3891



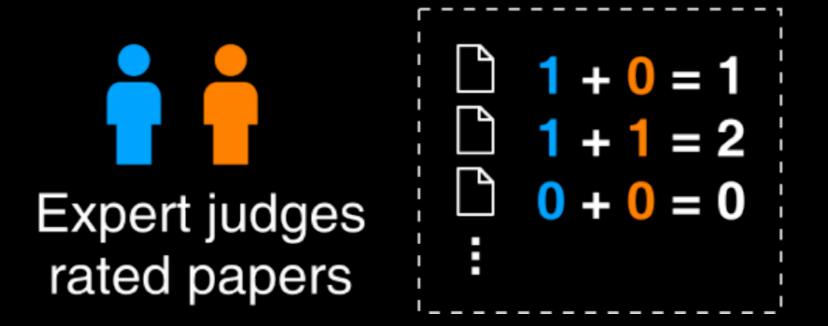


Between subjects design Participants: grad student or research staff

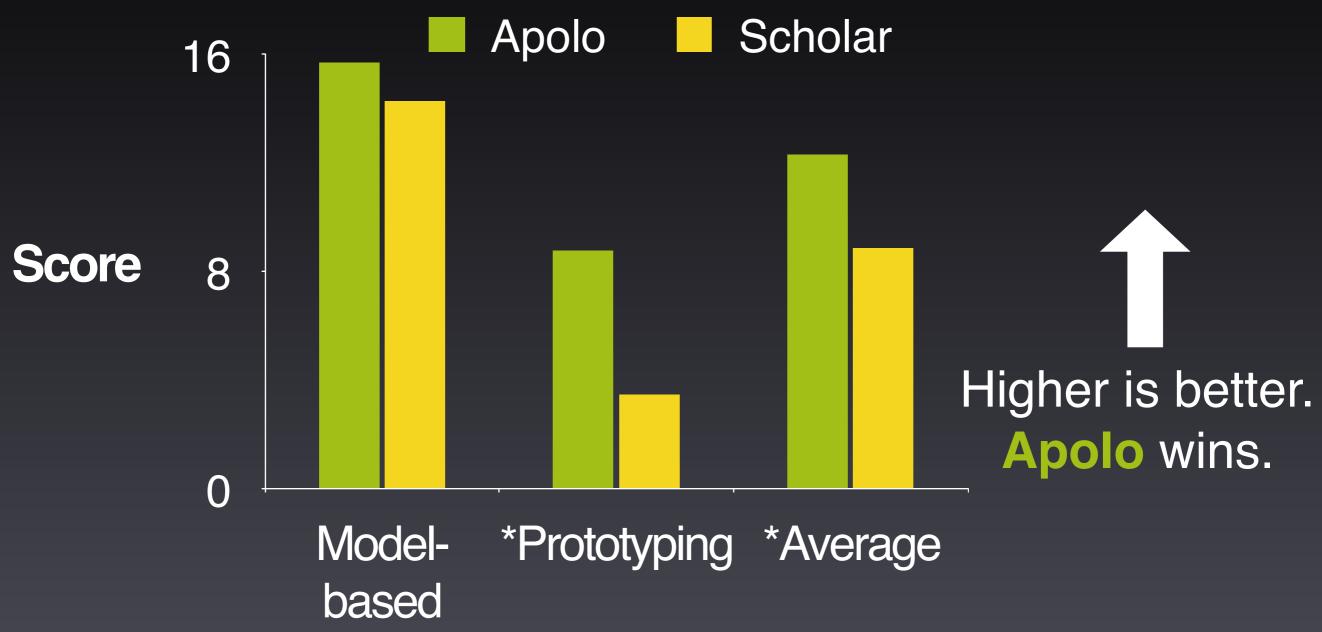




10 papers for each section



Judges' Scores



^{*} Statistically significant, by two-tailed t test, p < 0.05

Practitioners' guide to building (interactive) applications

What kinds of **prototypes**?

- Paper prototype, lo-fi prototype, high-fi prototype
 Important to involve REAL users as early as possible
 - Recruit your friends to try your tools
 - Lab study (controlled, as in Apolo)
 - Longitudinal study (usage over months)
 - Deploy it and see the world's reaction!
- To learn more:
 - CS 6750 Human-Computer Interaction
 - CS 6455 User Interface Design and Evaluation

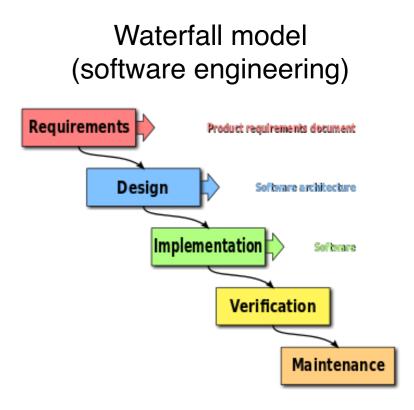
Practitioners' guide to building (interactive) applications

Think about scalability early

 Identify candidate scalable algorithms early on

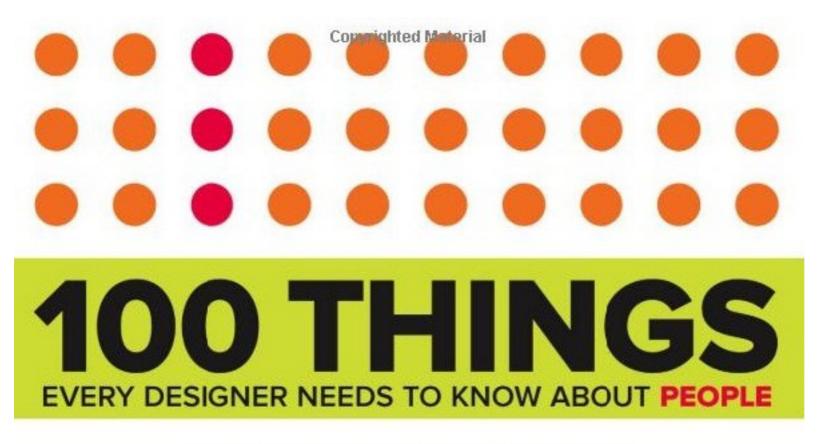
Use **iterative** design approach, as in Apolo and industry

- Why? It's hard to get it right the first time
- Create prototype, evaluate, modify prototype, evaluate, ...
- Quick evaluation helps you identify important fixes early — save you a lot of time overall



If you want to know more about people...

http://amzn.com/0321767535



SUSAN M. WEINSCHENK, Ph.D.

